

IST 707 Data Analytics HOMEWORK 4: Disease Prediction for Patients using LR, DT, ANN0, ANN1 and ANN2

Submitted by BHAVISH KUMAR on April 26th 2020

A. Executive Summary:

The purpose of this assignment is to use the Decision Tree, Logistic Regression classification models as well as the Artificial Neural Networks (Deep Learning) models that we have learned in the course to help solve the binary classification problem of predicting if a patient has the disease or not. The goal is to accurately predict if a given patient has the disease or not, thereby producing very good Accuracy, Recall and AUROC scores. Different Supervised Machine Learning algorithms were used to build models on the training dataset (dataset which contains the Target variable 'Disease' with its 2 classes) and these models were later used to make predictions on the new test dataset. The Machine Learning algorithms that were used to achieve this outcome were Decision Tree and Logistic Regression as well three Deep Learning algorithms with 0 hidden layers (ANN0), 1 hidden layer (ANN1) and 2 hidden layers (ANN2) were used. The goal of each model is to produce unbiased and low variance predictions which was achieved by extensive hyperparameter tuning done through a grid search. Several model evaluation techniques, which we will observe in the upcoming sections of the report were used to produce the best model.

B. INTRODUCTION:

In the below Analysis the following Learning Algorithms have been built: i. Logistic Regression: This is a regression based algorithm that converts the regression output to a probability ranging between 0 & 1 by using the sigmoid activation function. The Elastic Net Regularization parameters such as alpha and lambda can be used to control for overfitting.

ii. Decision Tree: This uses a tree-like model of decisions and their possible consequences. The tree contains a root node, internal nodes and leaf nodes along with splitting attributes that form the branches of the tree. The hyperparameters such as Max Depth, Minbucket, Minsplit etc. can be used to prune the tree in order to reduce variance without compromising on bias.

iii. Artificial Neural Network 0 layers (ANN0): This is a deep learning neural network model consisting of 0 hidden layers. This network consists of directly the output layer followed by the activation function to transform the output. The neural network uses Gradient Descent to identify the weights for the variables which will minimize the loss function.

iv. Artificial Neural Network 1 layer (ANN1): This is a feed forward Neural network containing 1 hidden layer followed by a relu activation function after the hidden layer and a an output layer followed by sigmoid activation function. The output layer consists of only 1 node and the sigmoid function in this case since this is a binary classification problem.

v. Artificial Neural Network 2 layers (ANN2): This is a feed forward Neural network containing 2 hidden layers followed by a relu activation function after each of the hidden layer and a an output layer followed by sigmoid activation function. The output layer consists of only 1 node and the sigmoid function in this case since this is a binary classification problem. The number of nodes and the layer dropout rate are parameters that can be tuned in order to reduce the variance as well as the bias. The batch size and

number of epochs are also two additional parameters which can be tuned to produce best results.

The following techniques were used for model performance evaluation:

K fold Cross Validation was done while training Logistic Regression & Decision Tree models on train data to control for overfitting, so that the model is built by ensuring that it performs well not just on train data but also on validation data which it has not seen before. Holdout method was also used where 70% of the data was used to train the model using K fold Cross Validation and the remaining 30% was used to measure the model performance on the data that it has not seen before. The model performance has been measured by using Accuracy, Precision, Recall, F1 score, ROC curve and Area Under ROC curve metrics.

C. Body of the Report:

Reading the Training data csv and storing it into a dataframe

```
setwd("D:/SYR ADS/Sem 2/IST_707_Data_Analytics/HW4")
getwd
```

```
## function ()
## .Internal(getwd())
## <bytecode: 0x000002478470eac8>
## <environment: namespace:base>
```

```
disease_prediction_training <- read.csv("Disease Prediction Training.csv")
```

VIEWING THE STRUCTURE and SUMMARY STATISTICS of the Data and checking for missing values

```
str(disease_prediction_training)
```

```
## 'data.frame':    49000 obs. of  12 variables:
## $ Age           : int  59 64 41 50 39 54 48 51 42 41 ...
## $ Gender        : Factor w/ 2 levels "female","male": 1 1 1 2 1 1 1 1 1 1 ...
## $ Height        : int  167 150 166 172 162 163 159 171 161 159 ...
## $ Weight        : num  88 71 83 110 61 61 89 71 72 43 ...
## $ High.Blood.Pressure: int  130 140 100 130 110 120 150 110 150 90 ...
## $ Low.Blood.Pressure : int  68 100 70 80 80 80 90 70 90 60 ...
## $ Cholesterol     : Factor w/ 3 levels "high","normal",...: 2 2 2 2 1 2 1 2 1 2 ...
## $ Glucose         : Factor w/ 3 levels "high","normal",...: 2 2 2 2 1 2 1 2 1 2 ...
## $ Smoke           : int  0 0 0 1 0 0 0 0 0 0 ...
## $ Alcohol         : int  0 0 1 0 0 0 0 0 0 0 ...
## $ Exercise        : int  1 0 1 1 1 1 1 1 1 1 ...
## $ Disease         : int  0 1 0 0 0 0 1 0 1 0 ...
```

```
summary(disease_prediction_training)
```

```
##      Age      Gender      Height      Weight      High.Blood.Pressure Low.Bloo
d.Pressure  Cholesterol      Glucose      Smoke
## Min.   :29.00  female:31863  Min.    : 55.0  Min.    : 10.00  Min.    : -150.0  Min.    :
0.00  high   : 6705  high    : 3627  Min.    :0.00000
## 1st Qu.:48.00  male  :17137  1st Qu.:159.0  1st Qu.: 65.00  1st Qu.: 120.0  1st Qu.:
80.00  normal :36676  normal :41652  1st Qu.:0.00000
## Median :53.00                      Median :165.0  Median : 72.00  Median : 120.0  Median :
80.00  too high: 5619  too high: 3721  Median :0.00000
## Mean   :52.85                      Mean   :164.4  Mean    : 74.19  Mean    : 128.7  Mean    :
96.92                      Mean    :0.08827
## 3rd Qu.:58.00                      3rd Qu.:170.0  3rd Qu.: 82.00  3rd Qu.: 140.0  3rd Qu.:
90.00                      3rd Qu.:0.00000
## Max.   :64.00                      Max.    :207.0  Max.    :200.00  Max.    :14020.0  Max.    :
11000.00                      Max.    :1.00000
##      Alcohol      Exercise      Disease
## Min.    :0.00000  Min.    :0.0000  Min.    :0.0
## 1st Qu.:0.00000  1st Qu.:1.0000  1st Qu.:0.0
## Median :0.00000  Median :1.0000  Median :0.0
## Mean    :0.05424  Mean    :0.8032  Mean    :0.5
## 3rd Qu.:0.00000  3rd Qu.:1.0000  3rd Qu.:1.0
## Max.    :1.00000  Max.    :1.0000  Max.    :1.0
```

SECTION 1: DATA PREPARATION & Exploratory Data Analysis

1. Identifying the Data Quality Issues:

As we can see from the above summary that the data has no missing values and hence NA imputation is not required

However, we can observe issues with 2 columns Low Blood Pressure and High Blood Pressure.

From the structure and summary of the data we can observe that the Min and Max values of the columns Low Blood Pressure and High Blood Pressure are not practically possible values and hence they are noise/outliers which need to be treated. Hence these columns need to be winsorized.

Winsorization is a data treatment process where the extreme outlier values are replaced with less extreme values which are practically possible

```
quantile(disease_prediction_training$Low.Blood.Pressure,c(0.001))
```

```
## 0.1%
## 45
```

```
quantile(disease_prediction_training$Low.Blood.Pressure,c(0.986))
```

```
## 98.6%
## 140
```

From the above 0.1 & 98.6 percentile values of low BP column we can observe that the possible values for the min & max of low BP (diastolic BP) fall in the range of 45 to 140 and hence any value that is less than 45 is replaced with 45 and any value greater than 140 is replaced with 140

```
disease_prediction_training$Low.Blood.Pressure[disease_prediction_training$Low.Blood.Pressure<quantile(disease_prediction_training$Low.Blood.Pressure,c(0.001))] <- quantile(disease_prediction_training$Low.Blood.Pressure,c(0.001))
```

```
disease_prediction_training$Low.Blood.Pressure[disease_prediction_training$Low.Blood.Pressure>quantile(disease_prediction_training$Low.Blood.Pressure,c(0.986))] <- quantile(disease_prediction_training$Low.Blood.Pressure,c(0.986))
```

Verifying that the Min & Max values of Low Blood Pressure (Diastolic BP) are in the correct practically permissible range and the outliers have been eliminated

```
summary(disease_prediction_training$Low.Blood.Pressure)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    45.00   80.00   80.00   82.15   90.00  140.00
```

The high Blood Pressure column also needs to winsorized to ensure that the values fall in the practically permissible range and outliers are eliminated

```
quantile(disease_prediction_training$High.Blood.Pressure,c(0.003))
```

```
## 0.3%
##    70
```

```
quantile(disease_prediction_training$High.Blood.Pressure,c(0.998))
```

```
## 99.8%
##   200
```

From the above 0.3 & 99.8 percentile values of High BP column we can observe that the possible values for the min & max of High BP (Systolic BP) fall in the range of 70 to 200 and hence any value that is less than 70 is replaced with 70 and any value greater than 200 is replaced with 200

```
disease_prediction_training$High.Blood.Pressure[disease_prediction_training$High.Blood.Pressure<quantile(disease_prediction_training$High.Blood.Pressure,c(0.003))] <- quantile(disease_prediction_training$High.Blood.Pressure,c(0.003))
```

```
disease_prediction_training$High.Blood.Pressure[disease_prediction_training$High.Blood.Pressure>quantile(disease_prediction_training$High.Blood.Pressure,c(0.998))] <- quantile(disease_prediction_training$High.Blood.Pressure,c(0.998))
```

Verifying that the Min & Max values of High Blood Pressure (Systolic BP) are in the correct practically permissible range and the outliers have been eliminated

```
summary(disease_prediction_training$High.Blood.Pressure)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      70.0   120.0   120.0   126.9   140.0   200.0
```

There are 245 instances where Low BP is > high BP even after winsorizing which needs to be treated by swapping the values

```
length(disease_prediction_training[disease_prediction_training$Low.Blood.Pressure>disease_prediction_training$High.Blood.Pressure,1])
```

```
## [1] 245
```

Swapping the values wherever Low BP > High BP which is not permissible

```
low_bp_values <- disease_prediction_training$Low.Blood.Pressure[disease_prediction_training$Low.Blood.Pressure>disease_prediction_training$High.Blood.Pressure]

high_bp_values <- disease_prediction_training$High.Blood.Pressure[disease_prediction_training$Low.Blood.Pressure>disease_prediction_training$High.Blood.Pressure]

disease_prediction_training$Low.Blood.Pressure[disease_prediction_training$Low.Blood.Pressure>disease_prediction_training$High.Blood.Pressure] <- high_bp_values

disease_prediction_training$High.Blood.Pressure[disease_prediction_training$Low.Blood.Pressure>disease_prediction_training$High.Blood.Pressure] <- low_bp_values
```

Verifying that there are no instances with low bp values > high bp values

```
length(disease_prediction_training[disease_prediction_training$Low.Blood.Pressure>disease_prediction_training$High.Blood.Pressure,1])
```

```
## [1] 0
```

The weight column has very low values, two of which are as low as 10Kg and 11Kg, which are practically very unlikely and hence they need to be winsorized

```
quantile(disease_prediction_training$Weight,c(0.0001))
```

```
##      0.01%
## 28.8999
```

Any value that is less than 0.01 percentile value, are replaced with the 0.01 percentile value = 28.9

```
disease_prediction_training$Weight[disease_prediction_training$Weight<quantile(disease_prediction_training$Weight,c(0.0001))] <- quantile(disease_prediction_training$Weight,c(0.0001))
```

Verifying that the Min & Max values of Weight are in the correct practically permissible range and the outliers have been eliminated

```
summary(disease_prediction_training$Weight)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  28.90   65.00   72.00   74.19   82.00  200.00
```

2. EXPLORATORY DATA ANALYSIS:

2.1. Bi Variate Analysis between Age and Disease column

Creating Age Groups column based on quartiles for EDA purpose

We assume that the number of people with the disease is higher for higher age groups, which we can verify by producing a bar graph

```
library(stringr)
disease_prediction_training$age_groups <- cut(disease_prediction_training$Age, breaks = c(quantile(disease_prediction_training$Age, probs = c(0,0.25,0.5,0.75,1))),
      labels = c(str_c(quantile(disease_prediction_training$Age,probs = 0),quantile(disease_prediction_training$Age,probs = 0.25),sep = " to "),str_c(quantile(disease_prediction_training$Age,probs = 0.25),quantile(disease_prediction_training$Age,probs = 0.5),sep = " to "),str_c(quantile(disease_prediction_training$Age,probs = 0.5),quantile(disease_prediction_training$Age,probs = 0.75),sep = " to "),str_c(quantile(disease_prediction_training$Age,probs = 0.75),quantile(disease_prediction_training$Age,probs = 1),sep = " to ")), right = FALSE, include.lowest=TRUE)
disease_prediction_training$age_groups<-as.factor(disease_prediction_training$age_groups)
#unique(disease_prediction_training$age_groups)
```

From the below bar group our assumption has been verified, as we can observe that the number of people with the disease increases as we go up the age groups and the older age groups have the highest number of patients with the disease

```
library(tidyverse)
```

```
## Warning: package 'tidyverse' was built under R version 3.6.2
```

```
## -- Attaching packages ----- tidyverse 1.3.0 --
```

```
## <U+2713> ggplot2 3.2.1      <U+2713> purrr  0.3.3
## <U+2713> tibble  2.1.3      <U+2713> dplyr  0.8.4
## <U+2713> tidyr   1.0.0      <U+2713> forcats 0.4.0
## <U+2713> readr   1.3.1
```


```
## Warning: package 'ggplot2' was built under R version 3.6.2
```

```
## Warning: package 'dplyr' was built under R version 3.6.2
```

```
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag() masks stats::lag()
```

```
library(ggplot2)
disease_count_by_ageGroups <- disease_prediction_training %>%
  group_by(age_groups)%>%
  summarise(sum(Disease))
colnames(disease_count_by_ageGroups) <- c('age_groups', 'NO_ppl_with_disease')
age_group_disease_plot <- ggplot(disease_count_by_ageGroups, aes(age_groups, NO_ppl_with_disease))
+geom_bar(stat = "identity")+
  xlab("Age Groups")+ ylab("Number of people with disease")+ ggtitle("Age Group VS co
unt of patients")
age_group_disease_plot
```

Age Group VS count of patients



2.2. Bi Variate Analysis between Height+Weight (BMI) and Disease column

Creating a new Body Mass Index (BMI) column by combining Height and Weight column, where $BMI = \frac{\text{Weight in Kg}}{(\text{Height in meters})^2}$

The BMI column can be used to classify the patients as Underweight, Healthy, Overweight and Obese. Underweight if BMI < 18.5; Healthy if BMI between 18.5 and 24.9; Overweight if BMI between 25 and 29.9; Obese if BMI greater than 30


```
disease_prediction_training$bmi <- disease_prediction_training$Weight/((disease_prediction_training$Height/100)*(disease_prediction_training$Height/100))
disease_prediction_training$bmi_groups <- cut(disease_prediction_training$bmi, breaks = c(0,18.5,24.9,29.9,Inf), labels = c('Underweight','Healthy','Overweight','Obese'))
```

We can observe that number of people with the disease is more for Overweight and Obese BMI groups in comparison to Healthy and Underweight BMI groups

```
disease_count_by_bmi_groups <- disease_prediction_training %>%
  group_by(bmi_groups)%>%
  summarise(sum(Disease))
colnames(disease_count_by_bmi_groups) <- c('bmi_groups', 'NO_ppl_with_disease')

bmi_disease_plot <- ggplot(disease_count_by_bmi_groups, aes(bmi_groups, NO_ppl_with_disease)) + geom_bar(stat = "identity") +
  xlab("bmi_groups") + ylab("Number of people with disease") + ggtitle("BMI groups VS count of patients")
bmi_disease_plot
```

BMI groups VS count of patients

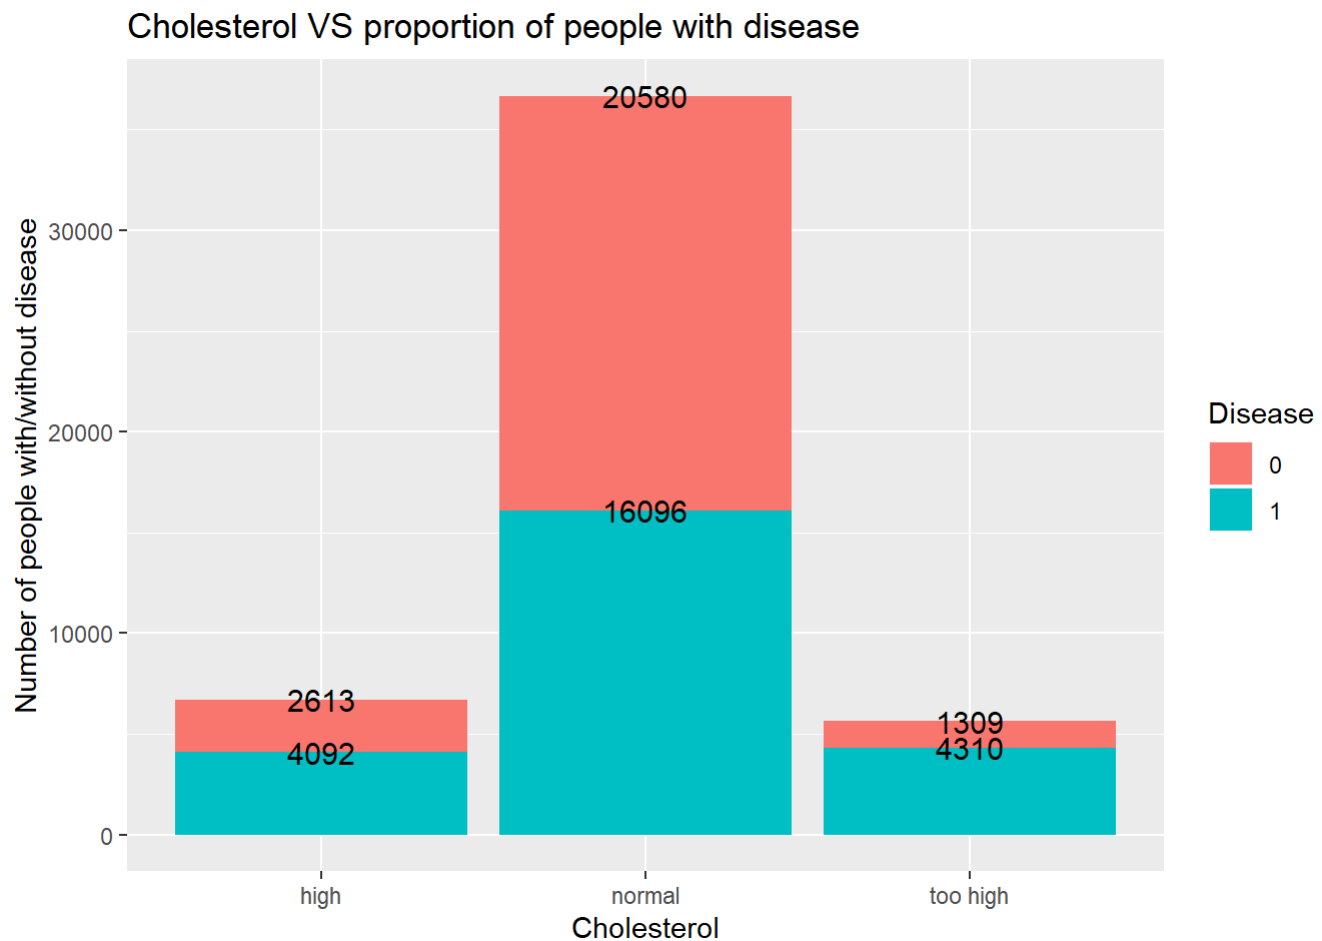


2.4. Bi Variate Analysis between Cholestrol and Disease column

We can observe that the percentage of people with the disease is much higher amongst 'high' and 'too high' Cholestrol groups of people


```
disease_count_by_cholesterol <- disease_prediction_training %>%
  group_by(Cholesterol,as.factor(Disease))%>%
  summarise(n())
colnames(disease_count_by_cholesterol) <- c('Cholesterol','Disease','NO_ppl')

Cholesterol_disease_plot <- ggplot(disease_count_by_cholesterol,aes(fill = Disease,x=Cholesterol,
y=NO_ppl))+geom_bar(position="stack",stat = "identity")+
  xlab("Cholesterol")+ ylab("Number of people with/without disease")+ ggtitle("Cholesterol VS proportion of people with disease")+ geom_text(aes(label = NO_ppl),position="stack",size = 4)
Cholesterol_disease_plot
```




2.5. Bi Variate Analysis between Glucose and Disease column

We can observe that the percentage of people with the disease is higher amongst 'high' and 'too high' Glucose groups of people

```
disease_count_by_Glucose <- disease_prediction_training %>%
  group_by(Glucose,as.factor(Disease))%>%
  summarise(n())
colnames(disease_count_by_Glucose) <- c('Glucose','Disease','NO_ppl')

Glucose_disease_plot <- ggplot(disease_count_by_Glucose,aes(fill = Disease,x=Glucose,y=NO_ppl))+
  geom_bar(position="stack",stat = "identity")+
  xlab("Glucose")+ ylab("Number of people with/without disease")+ ggtitle("Glucose VS
proportion of people with disease")+ geom_text(aes(label = NO_ppl),position="stack",size = 4)
Glucose_disease_plot
```



2.11. Uni Variate Analysis - Histogram of Numeric Variables to view the distribution of the numeric variables

The Weight column is mostly normally distributed with a slight right skew resulting in mean>median

The Low.Blood.Pressure column is mostly normally distributed with a slight right skew resulting in mean>median

The High.Blood.Pressure column is mostly normally distributed with a slight right skew resulting in mean>median

The Height column distribution is mostly normally distributed

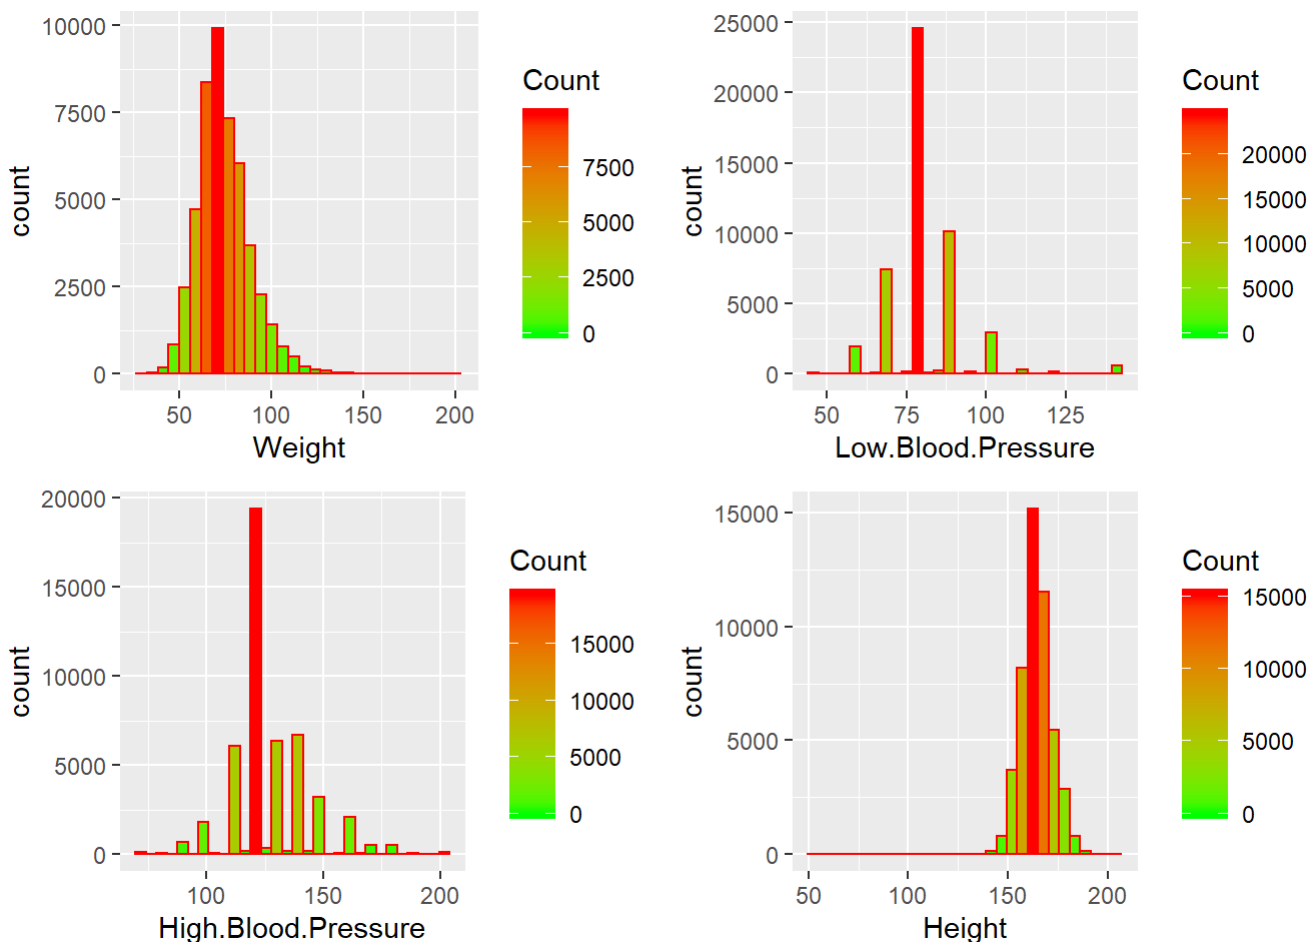
```
library(gridExtra)
```

```
## Warning: package 'gridExtra' was built under R version 3.6.3
```

```
##
## Attaching package: 'gridExtra'
```

```
## The following object is masked from 'package:dplyr':
##
##      combine
```

```
hist_plot <- function(data_in, i)
{
  data <- data.frame(x=data_in[[i]])
  p <- ggplot(data=data, aes(data[,1])) +
    geom_histogram(bins = 30, col="red", aes(fill=..count..)) +
    xlab(colnames(data_in)[i]) +
    scale_fill_gradient("Count", low="green", high="red")
  return (p)
}
mygrid2 <- list()
for (i in 1:length(disease_prediction_training[,c('Weight', 'Low.Blood.Pressure', 'High.Blood.Pressure', 'Height')])){
  myplot2 <- hist_plot(disease_prediction_training[,c('Weight', 'Low.Blood.Pressure', 'High.Blood.Pressure', 'Height')], i)
  mygrid2 <- c(mygrid2, list(myplot2))
}
do.call("grid.arrange", c(mygrid2, ncol=2))
```



3. DATA PREPARATION for Logistic & ANN :

For Logistic and ANN we require all the categorical variables to be converted to Numeric by performing one hot encoding and also all the numeric variables have to be normalized so that all the columns have values in the range of 0 to 1.

Creating Dummy Variables out of all categorical variables by performing one hot encoding and normalizing all other numeric columns using Min Max scaler

```
normalize <- function(x) {  
  return ((x - min(x)) / (max(x) - min(x)))  
}  
library(fastDummies)
```

```
## Warning: package 'fastDummies' was built under R version 3.6.2
```

```
disease_prediction_training_df <- disease_prediction_training  
disease_prediction_training_df <- fastDummies::dummy_cols(disease_prediction_training_df, select_columns=c('Gender', 'Cholesterol', 'Glucose'))  
  
disease_prediction_training_df$Age <- normalize(disease_prediction_training_df$Age)  
disease_prediction_training_df$Height <- normalize(disease_prediction_training_df$Height)  
disease_prediction_training_df$Weight <- normalize(disease_prediction_training_df$Weight)  
disease_prediction_training_df$bmi <- normalize(disease_prediction_training_df$bmi)  
disease_prediction_training_df$Low.Blood.Pressure <- normalize(disease_prediction_training_df$Low.Blood.Pressure)  
disease_prediction_training_df$High.Blood.Pressure <- normalize(disease_prediction_training_df$High.Blood.Pressure)
```

Getting rid of non numeric columns

```
disease_prediction_training_df$Gender <- NULL  
disease_prediction_training_df$Cholesterol <- NULL  
disease_prediction_training_df$Glucose <- NULL  
disease_prediction_training_df$age_groups <- NULL  
  
disease_prediction_training_df$bmi_groups <- NULL  
disease_prediction_training_df$Gender_male <- NULL  
disease_prediction_training_df$Disease <- as.factor(disease_prediction_training_df$Disease)  
  
i <- sapply(disease_prediction_training_df, is.integer)  
disease_prediction_training_df[,i] <- lapply(disease_prediction_training_df[,i], as.factor)
```

SECTION 2: BUILD, TUNE & EVALUATE various ML ALGORITHMS

2.1 Logistic Regression

Generating Training and Validation datasets for KNN algorithm with 70% & 30% splits respectively. (Hold one Out method to ensure that the model built doesn't overfit on train data by comparing the accuracies obtained on training data and validation data)

```
library(caret)
```

```
## Warning: package 'caret' was built under R version 3.6.2
```

```
## Loading required package: lattice
```

```
##  
## Attaching package: 'caret'
```

```
## The following object is masked from 'package:purrr':  
##  
## lift
```

```
set.seed(73)  
train_index <- createDataPartition(disease_prediction_training_df$Disease, p = 0.7, list = FALSE  
)  
  
disease_prediction_training_df_tr_split <- disease_prediction_training_df[train_index, ]  
  
disease_prediction_training_df_test_split <- disease_prediction_training_df[-train_index, ]  
  
disease_prediction_training_df_tr_split$Height <- NULL  
disease_prediction_training_df_tr_split$Weight <- NULL  
disease_prediction_training_df_test_split$Height <- NULL  
disease_prediction_training_df_test_split$Weight <- NULL
```

BUILDING BASELINE MODEL for Logistic Regression

```
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 3.6.3
```

```
## Loading required package: Matrix
```

```
##  
## Attaching package: 'Matrix'
```

```
## The following objects are masked from 'package:tidyr':  
##  
## expand, pack, unpack
```

```
## Loaded glmnet 3.0-2
```

```
baseline_model_logistic <- train(Disease ~ ., data = disease_prediction_training_df_tr_split, method = "glmnet", family = "binomial")
```

```
predict_disease_logistic <- predict(baseline_model_logistic, newdata = disease_prediction_training_df_test_split)
```

The baseline model produced the following accuracies on the training data for different values of elastic net regularization parameters

```
baseline_model_logistic
```

```
## glmnet
##
## 34301 samples
##    14 predictor
##    2 classes: '0', '1'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 34301, 34301, 34301, 34301, 34301, 34301, ...
## Resampling results across tuning parameters:
##
##  alpha  lambda      Accuracy  Kappa
##  0.10   0.0004282798  0.7298731  0.4596417
##  0.10   0.0042827984  0.7299367  0.4597687
##  0.10   0.0428279836  0.7289992  0.4578877
##  0.55   0.0004282798  0.7300351  0.4599655
##  0.55   0.0042827984  0.7299529  0.4597998
##  0.55   0.0428279836  0.7281104  0.4560975
##  1.00   0.0004282798  0.7300287  0.4599525
##  1.00   0.0042827984  0.7297086  0.4593087
##  1.00   0.0428279836  0.7248386  0.4495295
##
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were alpha = 0.55 and lambda = 0.0004282798.
```

The baseline model produced the below shown accuracy and sensitivity on validation data which it has not seen before.

```
confusionMatrix(predict_disease_logistic,disease_prediction_training_df_test_split$Disease, positive = "1")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 5797 2444
##           1 1553 4905
##
##           Accuracy : 0.7281
##           95% CI : (0.7208, 0.7353)
##           No Information Rate : 0.5
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.4561
##
##           Mcnemar's Test P-Value : < 2.2e-16
##
##           Sensitivity : 0.6674
##           Specificity : 0.7887
##           Pos Pred Value : 0.7595
##           Neg Pred Value : 0.7034
##           Prevalence : 0.5000
##           Detection Rate : 0.3337
##           Detection Prevalence : 0.4393
##           Balanced Accuracy : 0.7281
##
##           'Positive' Class : 1
##
```

Fine tuning the hyperparameter 'alpha' & 'lambda' for producing the best possible accuracy on validation data.

The Hyperparameter 'alpha' in Logistic algorithm denotes the balance between L1 & L2 regularization, whereas lambda is shrinkage parameter which denotes the strength of regularization alpha & lambda are the elastic net regularization parameters of a logistic regression model and are used to control for overfitting by reducing the coefficients of the features as these parameters are added to the Loss function in order to reduce the coefficients of the features which will in turn reduce overfitting Fine tuning the model to produce maximum accuracy and using 5 fold Cross Validation to train the model by controlling for overfitting

```
tuned_model_logistic <- train(Disease ~ ., data = disease_prediction_training_df_tr_split, metho
d = "glmnet",family = "binomial",
                             tuneLength = 15,
                             trControl = trainControl(method = "repeatedcv",
                                                         number = 5, repeats = 3))
```

From the below table summarizing the best performing models and their corresponding accuracies we can see the accuracies for different values of elastic net regularization parameters alpha and lambda

```
print(tuned_model_logistic)
```

```
## glmnet
##
## 34301 samples
##    14 predictor
##    2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 3 times)
## Summary of sample sizes: 27440, 27440, 27442, 27441, 27441, 27440, ...
## Resampling results across tuning parameters:
##
##    alpha      lambda      Accuracy      Kappa
##    0.1000000  0.0007616012  0.7303385  0.4606717
##    0.1000000  0.0013543398  0.7303385  0.4606717
##    0.1000000  0.0024083945  0.7303385  0.4606717
##    0.1000000  0.0042827984  0.7303676  0.4607300
##    0.1000000  0.0076160121  0.7300761  0.4601469
##    0.1000000  0.0135433976  0.7300371  0.4600690
##    0.1000000  0.0240839450  0.7299108  0.4598163
##    0.1000000  0.0428279836  0.7299302  0.4598551
##    0.1000000  0.0761601214  0.7296484  0.4592913
##    0.1000000  0.1354339757  0.7279575  0.4559089
##    0.1000000  0.2408394504  0.7271607  0.4543148
##    0.1642857  0.0007616012  0.7302413  0.4604773
##    0.1642857  0.0013543398  0.7302413  0.4604773
##    0.1642857  0.0024083945  0.7302413  0.4604773
##    0.1642857  0.0042827984  0.7303579  0.4607105
##    0.1642857  0.0076160121  0.7302801  0.4605550
##    0.1642857  0.0135433976  0.7298622  0.4597192
##    0.1642857  0.0240839450  0.7298331  0.4596608
##    0.1642857  0.0428279836  0.7302607  0.4605159
##    0.1642857  0.0761601214  0.7292306  0.4584553
##    0.1642857  0.1354339757  0.7282880  0.4565696
##    0.1642857  0.2408394504  0.7276952  0.4553833
##    0.2285714  0.0007616012  0.7302413  0.4604773
##    0.2285714  0.0013543398  0.7302413  0.4604773
##    0.2285714  0.0024083945  0.7303384  0.4606717
##    0.2285714  0.0042827984  0.7302899  0.4605745
##    0.2285714  0.0076160121  0.7302024  0.4603996
##    0.2285714  0.0135433976  0.7300080  0.4600107
##    0.2285714  0.0240839450  0.7298720  0.4597385
##    0.2285714  0.0428279836  0.7299691  0.4599326
##    0.2285714  0.0761601214  0.7286086  0.4572111
##    0.2285714  0.1354339757  0.7282393  0.4564721
##    0.2285714  0.2408394504  0.7279867  0.4559658
##    0.2928571  0.0007616012  0.7302899  0.4605745
##    0.2928571  0.0013543398  0.7302899  0.4605745
##    0.2928571  0.0024083945  0.7302899  0.4605745
##    0.2928571  0.0042827984  0.7302704  0.4605356
##    0.2928571  0.0076160121  0.7303093  0.4606133
##    0.2928571  0.0135433976  0.7300858  0.4601662
##    0.2928571  0.0240839450  0.7299886  0.4599717
##    0.2928571  0.0428279836  0.7299011  0.4597964
```


##	0.2928571	0.0761601214	0.7284920	0.4569777
##	0.2928571	0.1354339757	0.7281422	0.4562773
##	0.2928571	0.2408394504	0.7256057	0.4512029
##	0.3571429	0.0007616012	0.7302024	0.4603995
##	0.3571429	0.0013543398	0.7302024	0.4603995
##	0.3571429	0.0024083945	0.7301052	0.4602052
##	0.3571429	0.0042827984	0.7302316	0.4604579
##	0.3571429	0.0076160121	0.7303482	0.4606911
##	0.3571429	0.0135433976	0.7305037	0.4610020
##	0.3571429	0.0240839450	0.7301052	0.4602047
##	0.3571429	0.0428279836	0.7295707	0.4591354
##	0.3571429	0.0761601214	0.7282491	0.4564916
##	0.3571429	0.1354339757	0.7276563	0.4553051
##	0.3571429	0.2408394504	0.7189296	0.4378496
##	0.4214286	0.0007616012	0.7301830	0.4603607
##	0.4214286	0.0013543398	0.7301830	0.4603607
##	0.4214286	0.0024083945	0.7301149	0.4602246
##	0.4214286	0.0042827984	0.7303676	0.4607300
##	0.4214286	0.0076160121	0.7303773	0.4607493
##	0.4214286	0.0135433976	0.7304259	0.4608464
##	0.4214286	0.0240839450	0.7301343	0.4602630
##	0.4214286	0.0428279836	0.7289779	0.4579497
##	0.4214286	0.0761601214	0.7278410	0.4556751
##	0.4214286	0.1354339757	0.7272384	0.4544690
##	0.4214286	0.2408394504	0.7179383	0.4358680
##	0.4857143	0.0007616012	0.7301927	0.4603801
##	0.4857143	0.0013543398	0.7301927	0.4603801
##	0.4857143	0.0024083945	0.7302510	0.4604967
##	0.4857143	0.0042827984	0.7303676	0.4607300
##	0.4857143	0.0076160121	0.7303773	0.4607493
##	0.4857143	0.0135433976	0.7305328	0.4610601
##	0.4857143	0.0240839450	0.7303481	0.4606906
##	0.4857143	0.0428279836	0.7284532	0.4569001
##	0.4857143	0.0761601214	0.7276077	0.4552084
##	0.4857143	0.1354339757	0.7240801	0.4481515
##	0.4857143	0.2408394504	0.7178023	0.4355961
##	0.5500000	0.0007616012	0.7301732	0.4603412
##	0.5500000	0.0013543398	0.7301732	0.4603412
##	0.5500000	0.0024083945	0.7302024	0.4603995
##	0.5500000	0.0042827984	0.7304065	0.4608077
##	0.5500000	0.0076160121	0.7303967	0.4607881
##	0.5500000	0.0135433976	0.7303287	0.4606519
##	0.5500000	0.0240839450	0.7301149	0.4602239
##	0.5500000	0.0428279836	0.7278313	0.4556560
##	0.5500000	0.0761601214	0.7275494	0.4550915
##	0.5500000	0.1354339757	0.7192503	0.4384912
##	0.5500000	0.2408394504	0.7177342	0.4354600
##	0.6142857	0.0007616012	0.7301732	0.4603412
##	0.6142857	0.0013543398	0.7301635	0.4603218
##	0.6142857	0.0024083945	0.7303093	0.4606134
##	0.6142857	0.0042827984	0.7304551	0.4609048
##	0.6142857	0.0076160121	0.7304648	0.4609241
##	0.6142857	0.0135433976	0.7301149	0.4602242
##	0.6142857	0.0240839450	0.7300080	0.4600101

##	0.6142857	0.0428279836	0.7273939	0.4547813
##	0.6142857	0.0761601214	0.7270830	0.4541584
##	0.6142857	0.1354339757	0.7189684	0.4379281
##	0.6142857	0.2408394504	0.7177537	0.4354989
##	0.6785714	0.0007616012	0.7301732	0.4603412
##	0.6785714	0.0013543398	0.7301927	0.4603801
##	0.6785714	0.0024083945	0.7303482	0.4606911
##	0.6785714	0.0042827984	0.7302704	0.4605356
##	0.6785714	0.0076160121	0.7304842	0.4609630
##	0.6785714	0.0135433976	0.7300177	0.4600298
##	0.6785714	0.0240839450	0.7290848	0.4581635
##	0.6785714	0.0428279836	0.7273356	0.4546644
##	0.6785714	0.0761601214	0.7260820	0.4521562
##	0.6785714	0.1354339757	0.7179480	0.4358876
##	0.6785714	0.2408394504	0.7177537	0.4354989
##	0.7428571	0.0007616012	0.7301441	0.4602829
##	0.7428571	0.0013543398	0.7301635	0.4603218
##	0.7428571	0.0024083945	0.7304162	0.4608272
##	0.7428571	0.0042827984	0.7301732	0.4603411
##	0.7428571	0.0076160121	0.7304162	0.4608270
##	0.7428571	0.0135433976	0.7300080	0.4600103
##	0.7428571	0.0240839450	0.7282103	0.4564143
##	0.7428571	0.0428279836	0.7273939	0.4547808
##	0.7428571	0.0761601214	0.7243133	0.4486183
##	0.7428571	0.1354339757	0.7177440	0.4354795
##	0.7428571	0.2408394504	0.7177537	0.4354989
##	0.8071429	0.0007616012	0.7301538	0.4603023
##	0.8071429	0.0013543398	0.7301247	0.4602440
##	0.8071429	0.0024083945	0.7303871	0.4607688
##	0.8071429	0.0042827984	0.7301733	0.4603412
##	0.8071429	0.0076160121	0.7302996	0.4605937
##	0.8071429	0.0135433976	0.7299691	0.4599325
##	0.8071429	0.0240839450	0.7278216	0.4556368
##	0.8071429	0.0428279836	0.7273065	0.4546057
##	0.8071429	0.0761601214	0.7210773	0.4421456
##	0.8071429	0.1354339757	0.7177342	0.4354600
##	0.8071429	0.2408394504	0.7178217	0.4356349
##	0.8714286	0.0007616012	0.7301538	0.4603024
##	0.8714286	0.0013543398	0.7302218	0.4604384
##	0.8714286	0.0024083945	0.7304648	0.4609243
##	0.8714286	0.0042827984	0.7301830	0.4603606
##	0.8714286	0.0076160121	0.7301441	0.4602826
##	0.8714286	0.0135433976	0.7301440	0.4602823
##	0.8714286	0.0240839450	0.7275397	0.4550730
##	0.8714286	0.0428279836	0.7267914	0.4535756
##	0.8714286	0.0761601214	0.7183854	0.4367614
##	0.8714286	0.1354339757	0.7177342	0.4354600
##	0.8714286	0.2408394504	0.7179092	0.4358098
##	0.9357143	0.0007616012	0.7301635	0.4603218
##	0.9357143	0.0013543398	0.7302607	0.4605162
##	0.9357143	0.0024083945	0.7303676	0.4607299
##	0.9357143	0.0042827984	0.7302315	0.4604577
##	0.9357143	0.0076160121	0.7301149	0.4602243
##	0.9357143	0.0135433976	0.7299885	0.4599712

```
## 0.9357143 0.0240839450 0.7271413 0.4542759
## 0.9357143 0.0428279836 0.7261694 0.4523314
## 0.9357143 0.0761601214 0.7185311 0.4370532
## 0.9357143 0.1354339757 0.7177342 0.4354600
## 0.9357143 0.2408394504 0.5000437 0.0000000
## 1.0000000 0.0007616012 0.7301635 0.4603218
## 1.0000000 0.0013543398 0.7302704 0.4605356
## 1.0000000 0.0024083945 0.7304259 0.4608465
## 1.0000000 0.0042827984 0.7301149 0.4602245
## 1.0000000 0.0076160121 0.7299691 0.4599327
## 1.0000000 0.0135433976 0.7294929 0.4589799
## 1.0000000 0.0240839450 0.7273550 0.4547034
## 1.0000000 0.0428279836 0.7255572 0.4511068
## 1.0000000 0.0761601214 0.7183368 0.4366649
## 1.0000000 0.1354339757 0.7177342 0.4354600
## 1.0000000 0.2408394504 0.5000437 0.0000000
##
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were alpha = 0.4857143 and lambda = 0.0135434.
```

Obtaining Predicted values for disease on validation data using the fine tuned Logistic model

```
predict_disease_tuned_logistic <- predict(tuned_model_logistic, newdata = disease_prediction_training_df_test_split)
```

Evaluating Model Performance by calculating Accuracy Precision and Recall on the classification done on the Validation Data. (Hold one out method to control for overfitting)

After fine tuning the alpha and lambda elastic net regularization hyperparameters, the risk of overfitting can be avoided

Since the training and validation metric scores are in line, we can be assured that the model is producing low bias and low variance estimate

```
disease_prediction_training_df_test_split$predicted_disease_logistic <- predict_disease_tuned_logistic
conf_matrix <- data.frame(table(disease_prediction_training_df_test_split$Disease,disease_prediction_training_df_test_split$predicted_disease_logistic))
colnames(conf_matrix)<- c('Actual class','Predicted Class','Count')

Accuracy_DT <- sum(conf_matrix$Count[conf_matrix$`Actual class`==conf_matrix$`Predicted Class`])/sum(conf_matrix$Count)
Precision_DT <- conf_matrix$Count[conf_matrix$`Actual class`== 1 & conf_matrix$`Predicted Class`== 1]/sum(conf_matrix$Count[conf_matrix$`Predicted Class`==1])
Recall_DT <- conf_matrix$Count[conf_matrix$`Actual class`==1 & conf_matrix$`Predicted Class`==1]/sum(conf_matrix$Count[conf_matrix$`Actual class`==1])
F1_score_DT <- (2*Precision_DT*Recall_DT)/(Precision_DT+Recall_DT)

paste("Accuracy of best Logistic Algorithm in classifying Disease is :",Accuracy_DT)
```

```
## [1] "Accuracy of best Logistic Algorithm in classifying Disease is : 0.727396421525274"
```

```
paste("Precision of Logistic Algorithm in classifying Disease is :",Precision_DT)
```

```
## [1] "Precision of Logistic Algorithm in classifying Disease is : 0.76036148332814"
```

```
paste("Recall of Logistic Algorithm in classifying Disease is :",Recall_DT)
```

```
## [1] "Recall of Logistic Algorithm in classifying Disease is : 0.664035923254865"
```

```
paste("F1 Score of Logistic Algorithm in classifying Disease is :",F1_score_DT)
```

```
## [1] "F1 Score of Logistic Algorithm in classifying Disease is : 0.708941672114477"
```

ROC & AUROC for the best Logistic Regression Model

```
#install.packages("pROC")  
library(pROC)
```

```
## Warning: package 'pROC' was built under R version 3.6.2
```

```
## Type 'citation("pROC")' for a citation.
```

```
##  
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':  
##  
##      cov, smooth, var
```


Generated ROC curve and calculated Area Under Curve metric for the identified best performing Logistic Regression model

```
logistic_pred_disease <- predict(tuned_model_logistic, newdata = disease_prediction_training_df_  
test_split, na.action = na.omit, type= "prob")  
roc_curve <- roc(disease_prediction_training_df_test_split$Disease,logistic_pred_disease$`1`)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
plot(roc_curve)
```



Obtained the area under the ROC curve for the best Logistic model

```
paste("Area Under the ROC curve is :",auc(roc_curve))
```

```
## [1] "Area Under the ROC curve is : 0.789499122005585"
```

END OF BUILDING LOGISTIC Model (ML algorithm number 1)

2.2 Artificial Neural Network - 0 layers (ANN0)

Data Splitting

```
library(rsample)
```

```
## Warning: package 'rsample' was built under R version 3.6.3
```

```

train_test_split <- initial_split(disease_prediction_training_df, prop = 0.8)

disease_prediction_train_df_spit <- training(train_test_split)
disease_prediction_test_df_spit <- testing(train_test_split)

i <- sapply(disease_prediction_train_df_spit, is.factor)
disease_prediction_train_df_spit[,i] <- lapply(disease_prediction_train_df_spit[i], as.character)

j <- sapply(disease_prediction_train_df_spit, is.character)
disease_prediction_train_df_spit[,j] <- lapply(disease_prediction_train_df_spit[j], as.integer)

i <- sapply(disease_prediction_test_df_spit, is.factor)
disease_prediction_test_df_spit[,i] <- lapply(disease_prediction_test_df_spit[i], as.character)

j <- sapply(disease_prediction_test_df_spit, is.character)
disease_prediction_test_df_spit[,j] <- lapply(disease_prediction_test_df_spit[j], as.integer)

```

Data PreProcessing, creating rec_obj

step_center() to mean-center the data & **step_scale()** to scale the data

The last step is to prepare the recipe with the **prep()** function

```

#library(recipes)
disease_rec_obj <- recipes::recipe(Disease~.,data = disease_prediction_train_df_spit)%>%

  recipes::step_center(Age,Height,Weight,High.Blood.Pressure,Low.Blood.Pressure,bmi)%>%
  recipes::step_scale(Age,Height,Weight,High.Blood.Pressure,Low.Blood.Pressure,bmi)%>%
  recipes::prep(data = disease_prediction_train_df_spit)

```

Data PreProcessing: Creating Train and Test datasets and Train Test target variable Vectors

Applying the “recipe” to train and test data sets with the bake() function which processes the data following our recipe steps

```

x_train_disease_pred_tbl <- recipes::bake(disease_rec_obj, new_data = disease_prediction_train_df_spit)%>% select(-Disease)
x_test_disease_pred_tbl <- recipes::bake(disease_rec_obj, new_data = disease_prediction_test_df_spit)%>%select(-Disease)
y_train_vec <- ifelse(pull(disease_prediction_train_df_spit, Disease) == "1",1,0)
y_test_vec <- ifelse(pull(disease_prediction_test_df_spit, Disease) == "1",1,0)
str(x_train_disease_pred_tbl)

```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':   39201 obs. of  16 variables:
## $ Age                : num  0.909 1.649 -1.752 -0.422 -2.048 ...
## $ Height             : num  0.324 -1.742 0.203 0.932 -0.283 ...
## $ Weight             : num  0.967 -0.223 0.617 2.505 -0.922 ...
## $ High.Blood.Pressure : num  0.179 0.757 -1.553 0.179 -0.976 ...
## $ Low.Blood.Pressure  : num  -1.217 1.558 -1.043 -0.176 -0.176 ...
## $ Smoke              : int   0 0 0 1 0 0 0 0 0 0 ...
## $ Alcohol            : int   0 0 1 0 0 0 0 0 0 0 ...
## $ Exercise           : int   1 0 1 1 1 1 1 1 1 1 ...
## $ bmi                : num  0.653 0.653 0.418 1.574 -0.707 ...
## $ Gender_female      : int   1 1 1 0 1 1 1 1 1 0 ...
## $ Cholesterol_high    : int   0 0 0 0 1 0 1 1 0 0 ...
## $ Cholesterol_normal  : int   1 1 1 1 0 1 0 0 1 1 ...
## $ Cholesterol_too high: int   0 0 0 0 0 0 0 0 0 0 ...
## $ Glucose_high       : int   0 0 0 0 1 0 1 1 0 0 ...
## $ Glucose_normal     : int   1 1 1 1 0 1 0 0 1 1 ...
## $ Glucose_too high   : int   0 0 0 0 0 0 0 0 0 0 ...
```

```
#use_condaenv("r-tensorflow")
```

```
#install_tensorflow(package_url = "https://pypi.python.org/packages/b8/d6/af3d52dd52150ec4a6ceb7788bfeb2f62ecb6aa2d1172211c4db39b3#49a2/tensorflow-1.3.0rc0-cp27-cp27mu-manylinux1_x86_64.whl#md5=1cf77a2360ae2e38dd3578618eacc03b")
```

Building Deep Learning Models

Building ANN with 0 hidden layers(ANN0)

The ANN0 has no hidden layers and only has one output layer.

The output layer has one node as we are solving a binary classification problem and uses the sigmoid function to produce the output between 0 and 1

Using the “binary_crossentropy” loss function since this is a binary classification problem, and the goal of the 0 hidden layer ANN0 model is to minimize the binary_crossentropy loss through Gradient Descent Also writing a function below, which takes in the ‘optimizer’ as the input parameter and returns the 0 layer ANN architecture as output. We will compare the performance with both ADAM & SGD optimizers

```
library(keras)
```

```
## Warning: package 'keras' was built under R version 3.6.3
```

```
##
## Attaching package: 'keras'
```

```
## The following object is masked _by_ '.GlobalEnv':
##
##      normalize
```

```
library(tensorflow)
```

```
## Warning: package 'tensorflow' was built under R version 3.6.3
```

```
##  
## Attaching package: 'tensorflow'
```

```
## The following object is masked from 'package:caret':  
##  
##      train
```

```
model_keras_ann0 <- function(opt) {  
  
  model_keras_ann <- keras_model_sequential()  
model_keras_ann %>%  
  layer_dense(units = 1,  
    kernel_initializer = "uniform",  
    activation = "sigmoid") %>%  
    compile(optimizer = opt,  
      loss = "binary_crossentropy",  
      metrics = c("accuracy")  
  )  
  return (model_keras_ann)  
}
```


Fitting the Neural Network with 0 layers onto the train data using “ADAM” optimizer

```
model_ann0_fit <- fit(verbose=0,  
  object = model_keras_ann0("adam"),  
  x = as.matrix(x_train_disease_pred_tbl),  
  y = y_train_vec,  
  batch_size = 200,  
  epochs = 50,  
  validation_split = 0.30  
)
```

From the below plot we can observe how training & validation accuracies as well as losses vary with every epoch.

The point where the validation accuracy & loss curve begins to flatten will be the point where we can stop training the neural network

```
plot(model_ann0_fit)
```


Assesing the model performance using Stochastic Gradient Descent optimization algorithm

Fitting the Neural Network with 0 layers onto the train data using SGD optimizer

```
model_ann0_sgdfit <- fit(verbose=0,
  object = model_keras_ann0("SGD"),
  x = as.matrix(x_train_disease_pred_tbl),
  y = y_train_vec,
  batch_size = 200,
  epochs = 50,
  validation_split = 0.30
)
```

From the below plot we can observe how the training and validation accuracies & losses vary after every epoch

```
plot(model_ann0_sgdfit)
```



We can observe that there is not much significant difference in performance between SGD & ADAM optimizers (The training and validation accuracies are in line for both ADAM & SGD optimizers), hence going ahead with ADAM, as it is marginally better than SGD

```
paste("MODEL PERFORMANCE with SGD optimizer",model_ann0_sgdfit)
```

```
## [1] "MODEL PERFORMANCE with SGD optimizer list(batch_size = 200, epochs = 50, steps = 138, samples = 27440, verbose = 0, do_validation = TRUE, metrics = c(\"loss\", \"accuracy\", \"val_loss\", \"val_accuracy\"))"
## [2] "MODEL PERFORMANCE with SGD optimizer list(loss = c(0.640947024172319, 0.599813293872005, 0.584758857075049, 0.577840238809586, 0.574053966529863, 0.571720123986113, 0.570171061392776, 0.569065503928126, 0.568234582814461, 0.56756776862346, 0.567015108665989, 0.566544360565375, 0.56614710729428, 0.565805925746006, 0.565489257737429, 0.565212579101932, 0.564960707684995, 0.564740976508783, 0.564531860693898, 0.564346161597672, 0.564172875038389, 0.564027532804812, 0.56387802689138, 0.563746517062535, 0.563627377586879, 0.563521958605193, 0.563422465550309, 0.56330646115559, 0.5632355204911, 0.563163639418238, 0.563078774828953, 0.563013483023505, 0.562948705242952, 0.562881647063414, 0.562832106958325, 0.562773106092962, 0.562723933434, 0.562667745847049, 0.562639228282795, 0.56258883186054, 0.562546493013121, 0.562505133471753, 0.562465838588362, 0.562431134392152, 0.562403600261093, 0.562366940624165, 0.562336855614151, 0.562304418783841, 0.562279450215682, 0.562251160817313), accuracy = c(0.682142853736877, 0.711880445480347, 0.717201173305511, 0.718913972377777, 0.720590353012085, 0.721829473972321, 0.723287165164948, 0.72434401512146, 0.724453330039978, 0.724562704563141, 0.724963545799255, 0.726749241352081, 0.726567029953003, 0.727150142192841, 0.727733254432678, 0.727733254432678, 0.728097677230835, 0.728243410587311, 0.728680729866028, 0.728607892990112, 0.728534996509552, 0.729118049144745, 0.729008734226227, 0.729409635066986, 0.729190945625305, 0.729008734226227, 0.729118049144745, 0.729336738586426, 0.729336738586426, 0.729774057865143, 0.729081630706787, 0.729190945625305, 0.728972315788269, 0.729263842105865, 0.729081630706787, 0.729227423667908, 0.729482531547546, 0.729336738586426, 0.729045212268829, 0.729154527187347, 0.729482531547546, 0.729446053504944, 0.729591846466064, 0.729409635066986, 0.729555368423462, 0.729336738586426, 0.729701161384583, 0.729409635066986, 0.729446053504944, 0.729518949985504), val_loss = c(0.612229292504835, 0.588854970995661, 0.578784378630482, 0.573572448114688, 0.570435131698956, 0.568446560626066, 0.567016355197311, 0.565976066883704, 0.565215583196678, 0.564549375959891, 0.564047064077549, 0.563574550103375, 0.563246395558324, 0.562881302416968, 0.562587346342204, 0.562312871711857, 0.562011445504185, 0.561788109484887, 0.561617523870159, 0.561437856055332, 0.561239840906095, 0.561090232340538, 0.560927261235651, 0.560827251458166, 0.560702667340931, 0.560568217502019, 0.560467471213641, 0.560338970856869, 0.560245324687359, 0.560191500669393, 0.56010092972655, 0.560016281106482, 0.559918271001226, 0.559872094622559, 0.559790839559822, 0.55975428391107, 0.559707875757961, 0.559715771328459, 0.559623364404436, 0.559570173297485, 0.559484750131251, 0.559441265393131, 0.559375805870407, 0.559373188820555, 0.559350304985298, 0.559279032436097, 0.559225316756579, 0.559185315922869, 0.559161047593498, 0.559160545360636), val_accuracy = c(0.707167744636536, 0.714054942131042, 0.715500354766846, 0.718221247196198, 0.71881639957428, 0.718476295471191, 0.719241559505463, 0.719581663608551, 0.718986451625824, 0.719496667385101, 0.719326615333557, 0.719751715660095, 0.720601975917816, 0.721282184123993, 0.721452236175537, 0.72187739610672, 0.722812712192535, 0.723067760467529, 0.723577916622162, 0.723747968673706, 0.723577916622162, 0.723833024501801, 0.723747968673706, 0.72391802072525, 0.723747968673706, 0.724343180656433, 0.723662972450256, 0.723577916622162, 0.723833024501801, 0.72391802072525, 0.724173128604889, 0.724428176879883, 0.72502338886261, 0.725363492965698, 0.725363492965698, 0.725278437137604, 0.724938333034515, 0.725193440914154, 0.725363492965698, 0.725703597068787, 0.725618541240692, 0.725533545017242, 0.725703597068787, 0.725193440914154, 0.724768280982971, 0.725448489189148, 0.726128756999969, 0.726383805274963, 0.726383805274963, 0.725703597068787))"
```

```
paste("MODEL PERFORMANCE with ADAM optimizer",model_ann0_fit)
```

```
## [1] "MODEL PERFORMANCE with ADAM optimizer list(batch_size = 200, epochs = 50, steps = 138, samples = 27440, verbose = 0, do_validation = TRUE, metrics = c("\loss", "\accuracy", "\val_loss", "\val_accuracy"))"
## [2] "MODEL PERFORMANCE with ADAM optimizer list(loss = c(0.665677072789857, 0.617283680956843, 0.594813277961214, 0.582656037181877, 0.575694208459673, 0.571612087005081, 0.569128292859817, 0.567451958232301, 0.566260115528593, 0.565310904739897, 0.564610833198962, 0.564062515344286, 0.563480261333135, 0.563111475535801, 0.56277735974455, 0.562557864206525, 0.562339211214041, 0.562141366802569, 0.562010313392381, 0.561904796054342, 0.561856921968585, 0.561794181259311, 0.561744006095405, 0.561653072476039, 0.561666187452853, 0.561695466409967, \n0.561630451210039, 0.561633256060389, 0.56165993183243, 0.561653129951947, 0.561603711793096, 0.561631871208158, 0.561612466551125, 0.561607389642963, 0.561605033174201, 0.56162929665243, 0.56161719834318, 0.561662558868049, 0.561614842569515, 0.561574467286772, 0.561626436411466, 0.561588743554955, 0.561599448888017, 0.561605889838917, 0.561593451062027, 0.561571806278243, 0.561620656107675, 0.561626491411087, 0.56161304096265, 0.561575083882051), accuracy = c(0.613265335559845, 0.690051019191742, \n0.704409599304199, 0.712463557720184, 0.716618061065674, 0.717857122421265, 0.721137046813965, 0.722266793251038, 0.723615169525146, 0.725218653678894, 0.725437343120575, 0.726895034313202, 0.727878987789154, 0.728389203548431, 0.729154527187347, 0.729227423667908, 0.729774057865143, 0.729883372783661, 0.729883372783661, 0.730976700782776, 0.730575799942017, 0.730612218379974, 0.7301384806633, 0.730903804302216, 0.730721592903137, 0.730065584182739, 0.730502903461456, 0.730393588542938, 0.730721592903137, \n0.730247795581818, 0.73035717010498, 0.730430006980896, 0.730758011341095, 0.730648696422577, 0.730502903461456, 0.730466485023499, 0.730612218379974, 0.730830907821655, 0.730320692062378, 0.730758011341095, 0.730247795581818, 0.73028427362442, 0.73035717010498, 0.730174899101257, 0.730539381504059, 0.730685114860535, 0.730721592903137, 0.730502903461456, 0.730393588542938, 0.730102062225342), val_loss = c(0.634121375311578, 0.602007217157273, 0.585738700705821, 0.57654950973706, 0.570973160350802, \n0.567562223393012, 0.565183649571126, 0.563812055693994, 0.56250044486827, 0.561842776993973, 0.561002788880419, 0.560314798859874, 0.559844208350545, 0.559505692946141, 0.559145234117864, 0.558944688936184, 0.558900889465673, 0.558794134080374, 0.558460820056723, 0.558613301109025, 0.55836958699603, 0.558498907049703, 0.558266419561002, 0.558392389327853, 0.558095124077811, 0.55819083919899, 0.558292590728494, 0.558108621025337, 0.558129865329957, 0.558086636261859, 0.558028559575893, 0.558345934706332, \n0.558134651647296, 0.558107615961589, 0.558085675041303, 0.558132083153869, 0.557981482615186, 0.558145254067383, 0.558079716908102, 0.558394433831532, 0.558133697273597, 0.558105362104261, 0.558144487944218, 0.55803299702792, 0.55810127877812, 0.557918511420405, 0.557974537647801, 0.558001543898615, 0.558023436424653, 0.557969523509053), val_accuracy = c(0.674857556819916, 0.701470971107483, 0.709633529186249, 0.716180622577667, 0.71881639957428, 0.721452236175537, 0.723662972450256, 0.722727656364441, \n0.72298276424408, 0.723492920398712, 0.724173128604889, 0.724088072776794, 0.723833024501801, 0.724513232707977, 0.725703597068787, 0.726213753223419, 0.72510838508606, 0.726298809051514, 0.727319121360779, 0.726638913154602, 0.727234065532684, 0.726723909378052, 0.727744221687317, 0.727234065532684, 0.727999329566956, 0.727744221687317, 0.72706401348114, 0.727744221687317, 0.727149069309235, 0.727914273738861, 0.728339433670044, 0.726808965206146, 0.727489173412323, 0.727659225463867, 0.728339433670044, \n0.728254377841949, 0.728084325790405, 0.72697901725769, 0.727489173412323, 0.72697901725769, 0.727999329566956, 0.72706401348114, 0.728339433670044, 0.728254377841949, 0.727829277515411, 0.72697901725769, 0.728084325790405, 0.727914273738861)))"
```

Tuning the number of epochs and batch size parameters by calling the `model_keras_ann0` function written above which returns a 0 layer ANN0 network in every iteration of the loop. Seeing which batch size & epochs combination produces the best validation accuracy using ADAM optimizer network

The number of epochs and batch size should be chosen in a way that both training and validation accuracies are high thus ensuring that the model doesn't overfit by producing a high training accuracy and low validation accuracy

```
batch_size = c(50,100,150,200)
epochs = c(10,15,20,40,60,100)
model_ann0_tuned_fit <- c(0)
for (i in batch_size) {
  for (j in epochs) {

    model_fit <- fit(verbose=0,
object = model_keras_ann0('adam'),
x = as.matrix(x_train_disease_pred_tbl),
y = y_train_vec,
batch_size = i,
epochs = j,
validation_split = 0.30
)
    model_ann0_tuned_fit <- c(model_ann0_tuned_fit,model_fit)
  }
}
```

Storing the training and validation accuracies for different batch size and epoch values to compare the accuracies for each combination

```
batch_sizes <- c()
epochs_run <- c()
validation_accuracies <- c()
training_accuracies <- c()
for (i in seq(2,48,2))
{
  batch_sizes <- c(batch_sizes,model_ann0_tuned_fit[i]$params$batch_size)
  epochs_run <- c(epochs_run,model_ann0_tuned_fit[i]$params$epochs)
}

for (i in seq(3,49,2)) {
  validation_accuracies <- c(validation_accuracies,tail(model_ann0_tuned_fit[i]$metrics$val_accuracy,n=1))
  training_accuracies <- c(training_accuracies, tail(model_ann0_tuned_fit[i]$metrics$accuracy,n=1))
}
```

Table summarizing the model performances and their corresponding batch size and epoch hyperparameters

The batch_size values, epochs values and their corresponding training accuracy as well as validation accuracy can be identified from the below table

```
performance_summary_df <- data.frame(batch_sizes,epochs_run,training_accuracies,validation_accuracies)
performance_summary_df
```

##	batch_sizes	epochs_run	training_accuracies	validation_accuracies
## 1	50	10	0.7307216	0.7274041
## 2	50	15	0.7315598	0.7277442
## 3	50	20	0.7305394	0.7279993
## 4	50	40	0.7302114	0.7264689
## 5	50	60	0.7305029	0.7273191
## 6	50	100	0.7304665	0.7271491
## 7	100	10	0.7301749	0.7265539
## 8	100	15	0.7310131	0.7266389
## 9	100	20	0.7302843	0.7278293
## 10	100	40	0.7316326	0.7266389
## 11	100	60	0.7315962	0.7281694
## 12	100	100	0.7303572	0.7275742
## 13	150	10	0.7277697	0.7241731
## 14	150	15	0.7305029	0.7280843
## 15	150	20	0.7309402	0.7283394
## 16	150	40	0.7308673	0.7275742
## 17	150	60	0.7305758	0.7268090
## 18	150	100	0.7302843	0.7269790
## 19	200	10	0.7247449	0.7240881
## 20	200	15	0.7293732	0.7267239
## 21	200	20	0.7312317	0.7275742
## 22	200	40	0.7303936	0.7269790
## 23	200	60	0.7304665	0.7266389
## 24	200	100	0.7311224	0.7271491

Identifying the number of Epochs and Batch Size which produced the maximum validation accuracy.

```
performance_summary_df$epochs_run[which.max(performance_summary_df$validation_accuracies)]
```

```
## [1] 20
```

```
performance_summary_df$batch_sizes[which.max(performance_summary_df$validation_accuracies)]
```

```
## [1] 150
```

Hence going ahead with the batch size and epochs that produced the maximum validation accuracy to fit the final ANNO deep learning model. When the validation accuracy is also as high as training accuracy, we can be assured that the model will produce low bias and low variance estimates.

```

dl_model_keras_ann0 <- model_keras_ann0('adam')
model_ann0_fit <- fit(verbose=0,
  object = dl_model_keras_ann0,
  x = as.matrix(x_train_disease_pred_tbl),
  y = y_train_vec,
  batch_size = performance_summary_df$batch_sizes[which.max(performance_summary_df$validation_accuracies)],
  epochs = performance_summary_df$epochs_run[which.max(performance_summary_df$validation_accuracies)],
  validation_split = 0.30
)

```

Making Predictions on the split Test Data using the final ANN0 deep learning model built above using the tibble function and yardstick library

Accuracy and Recall are the 2 most important metrics because Recall helps in minimizing the False Negatives which are highly undesirable when doing disease prediction

```

yhat_keras_class_vec <- predict_classes(object = dl_model_keras_ann0, x = as.matrix(x_test_disease_pred_tbl)) %>% as.vector()
yhat_keras_prob_vec <- predict_proba(object = dl_model_keras_ann0, x = as.matrix(x_test_disease_pred_tbl)) %>% as.vector()
estimates_keras_tbl <- tibble::tibble(
  truth = as.factor(y_test_vec),
  estimate = as.factor(yhat_keras_class_vec),
  class_prob = yhat_keras_prob_vec
)
#estimates_keras_tbl

```

Evaluating the ANN0 model performance on the test dataset (Hold One out method), using the same set of metrics: Accuracy, Precision, Recall, F1 score and AUC score to evaluate the performance of the ANN0 Deep Learning model

Since the accuracy on the testing data is completely inline with the training accuracy we can be assured that the model is producing low bias and low variance estimate

```
library(yardstick)
```

```
## Warning: package 'yardstick' was built under R version 3.6.3
```

```
## For binary classification, the first factor level is assumed to be the event.
## Set the global option `yardstick.event_first` to `FALSE` to change this.
```

```
##
## Attaching package: 'yardstick'
```

```
## The following object is masked from 'package:keras':
##
##   get_weights
```

```
## The following objects are masked from 'package:caret':
##
##   precision, recall, sensitivity, specificity
```

```
## The following object is masked from 'package:readr':
##
##   spec
```

```
options(yardstick.event_first = F)
ann0_accuracy <- estimates_keras_tbl %>% metrics(truth, estimate)
ACCURACY<- ann0_accuracy$.estimate[1]

ann0_auc <-estimates_keras_tbl %>% roc_auc(truth, class_prob)
AUC<-ann0_auc$.estimate[1]

ann0_precision <-estimates_keras_tbl %>% precision(truth, estimate)
PRECISION <- ann0_precision$.estimate[1]

ann0_recall <-estimates_keras_tbl %>% recall(truth, estimate)
RECALL<- ann0_recall$.estimate[1]

ann0_f1 <-estimates_keras_tbl %>% f_meas(truth, estimate)
F1 <- ann0_f1$.estimate[1]

data.frame(ACCURACY,AUC,PRECISION,RECALL,F1)
```

```
##   ACCURACY      AUC PRECISION  RECALL      F1
## 1 0.7272171 0.7933942 0.7522584 0.6788262 0.7136583
```

```
#install.packages("pROC")
library(pROC)
```


Generated ROC curve and calculated Area Under Curve metric for the identified best performing ANN0 DL model

```
roc_curve <- roc(y_test_vec,yhat_keras_prob_vec)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
plot(roc_curve)
```

Obtained the area under the ROC curve for the best ANN0 model

```
paste("Area Under the ROC curve is :",auc(roc_curve))
```

```
## [1] "Area Under the ROC curve is : 0.793394150829301"
```

END OF BUILDING ANN0 Deep Learning Model (algorithm number 2)

Comparing preformances of Linear SVM, Logistic Regression and ANN0 (Perceptron with 0 hidden layers)

```
## LOGISTIC REGRESSION
disease_prediction_training_df_test_split$predicted_disease_logistic <- predict_disease_tuned_logistic
conf_matrix <- data.frame(table(disease_prediction_training_df_test_split$Disease,disease_prediction_training_df_test_split$predicted_disease_logistic))
colnames(conf_matrix)<- c('Actual class','Predicted Class','Count')
Accuracy_Logistic <- sum(conf_matrix$Count[conf_matrix$`Actual class`==conf_matrix$`Predicted Class`])/sum(conf_matrix$Count)
Precision_Logistic <- conf_matrix$Count[conf_matrix$`Actual class`== 1 & conf_matrix$`Predicted Class`== 1]/sum(conf_matrix$Count[conf_matrix$`Predicted Class`==1])
Recall_Logistic <- conf_matrix$Count[conf_matrix$`Actual class`==1 & conf_matrix$`Predicted Class`==1]/sum(conf_matrix$Count[conf_matrix$`Actual class`==1])
F1_score_Logistic <- (2*Precision_DT*Recall_DT)/(Precision_DT+Recall_DT)
logistic_pred_disease <- predict(tuned_model_logistic, newdata = disease_prediction_training_df_test_split, na.action = na.omit, type = "prob")
roc_curve_logistic <- roc(disease_prediction_training_df_test_split$Disease,logistic_pred_disease$`1`)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
auc_logistic <- auc(roc_curve_logistic)
```

```
## PERFORMANCE COMPARISON
ALGORITHMS <- c('LOGISTIC REGRESSION','ANN0','LINEAR SVM')
ACCURACIES <- c(Accuracy_Logistic,ACCURACY,0.7288)
PRECISIONS <- c(Precision_Logistic,PRECISION,0.7988)
RECALLs <- c(Recall_Logistic,RECALL,0.61178)
F1_scores <- c(F1_score_Logistic,F1,0.72558)

data.frame(ALGORITHMS,ACCURACIES,PRECISIONs,RECALLs,F1_scores)
```

```
##          ALGORITHMS  ACCURACIES  PRECISIONs  RECALLs  F1_scores
## 1 LOGISTIC REGRESSION  0.7273964  0.7603615  0.6640359  0.7089417
## 2              ANN0    0.7272171  0.7522584  0.6788262  0.7136583
## 3          LINEAR SVM  0.7288000  0.7988000  0.6117800  0.7255800
```

Discussion on comparison of Linear SVM, Logistic Regression and ANN0 (Perceptron with 0 hidden layers)

The performances of Logistic Regression, Linear SVM and single layer perceptron (ANN0) are very much comparable because all the algorithms work on the concept of dividing the data points by constructing a linear hyperplane that separates the data into 2 classes. The single layer perceptron is also a linear combination of predictor variables (weighted sum of predictors) just like logistic regression, because of which both these algorithms construct a linear hyperplane separating the data points. Also linear SVM works on the exact same principle of constructing a linear hyperplane separating the data points, similar to logistic regression and ANN0. Moreover, the method of obtaining target variable probabilities is same in

Logistic Regression and ANN0, both these methods start of by taking a weighted sum (linear combination) of predictor variables and then apply the sigmoid function to suppress the output between 0 and 1. Hence all the 3 algorithms have similar performance.

2.3 Artificial Neural Network - 1 hidden layer (ANN1)

Using the same X train, X test dataset tables and Y train, Y test vectors

Building neural network architecture with 1 hidden layer (ANN1)

The number of nodes in the first hidden layer should be equal to the number of input attributes in the dataset

The Dropout layers are used to control overfitting by eliminating weights below a cutoff threshold in order to prevent low weights from overfitting the layers. Here we are removing weights below 10%.

Using the relu activation function after the hidden layer as relu produces 0 as output for all negative inputs and the input value itself is produced as output if the input is > 0. Hence, the relu activation function is more suitable for taking partial derivatives. The Sigmoid function is being used after the output layer to produce output between 0 and 1.

Using the binary_crossentropy loss function since it is a binary classification problem

Written a function below, which returns a Neural Network with one hidden layer by taking in the optimizer as input parameter

```
library(keras)
library(tensorflow)

model_keras_ann1 <- function(opt) {

model_keras_ann <- keras_model_sequential()
model_keras_ann %>%
  ## hidden layer 1
  layer_dense(units = 16,
    kernel_initializer = "uniform",
    activation = "relu",
    input_shape = ncol(x_train_disease_pred_tbl)) %>%
  layer_dropout(rate = 0.1)%>%
  ## Output Layer
  layer_dense(units = 1,
    kernel_initializer = "uniform",
    activation = "sigmoid") %>%
  compile(optimizer = "ADAM",
    loss = "binary_crossentropy",
    metrics = c("accuracy"))
)

return(model_keras_ann)

}
```

Fitting the Neural Network with 1 hidden layer onto the train data


```

model_ann1_fit <- fit(verbose=0,
  object = model_keras_ann1('adam'),
  x = as.matrix(x_train_disease_pred_tbl),
  y = y_train_vec,
  batch_size = 50,
  epochs = 15,
  validation_split = 0.30
)

```

From the below plot also we can identify the suitable number of epochs based on the point where the curve begins to flatten

```
plot(model_ann1_fit)
```



Tuning the number of epochs and batch size parameters to see which combination produces the best validation accuracy using ADAM optimizer network by using the ANN1 function written above which return a DL neural network with 1 hidden layer in every iteration of the loop, by taking the optimizer as the function input parameter.

The number of epochs & batch size should be chosen in a way that both training and validation accuracies are high thus ensuring that the model doesn't overfit by producing a high training accuracy and low validation accuracy

```

batch_size = c(50,100,150,200)
epochs = c(10,15,20,40,60,100)
model_ann1_tuned_fit <- c(0)
for (i in batch_size) {
  for (j in epochs) {

    model_fit <- fit(verbose=0,
object = model_keras_ann1('adam'),
x = as.matrix(x_train_disease_pred_tbl),
y = y_train_vec,
batch_size = i,
epochs = j,
validation_split = 0.30
)
    model_ann1_tuned_fit <- c(model_ann1_tuned_fit,model_fit)
  }
}

```

Storing the training and validation accuracies for different batch size and epoch values to compare the accuracies for each combination

```

batch_sizes <- c()
epochs_run <- c()
validation_accuracies <- c()
training_accuracies <- c()
for (i in seq(2,48,2))
{
  batch_sizes <- c(batch_sizes,model_ann1_tuned_fit[i]$params$batch_size)
  epochs_run <- c(epochs_run,model_ann1_tuned_fit[i]$params$epochs)
}

for (i in seq(3,49,2)) {
  validation_accuracies <- c(validation_accuracies,tail(model_ann1_tuned_fit[i]$metrics$val_accuracy,n=1))
  training_accuracies <- c(training_accuracies, tail(model_ann1_tuned_fit[i]$metrics$accuracy,n=1))
}

```

Table summarizing the model performances and their corresponding batch size and epoch hyperparameters

The batch_size value & epochs values with their corresponding training accuracy and validation accuracy can be seen in the below table

```

performance_summary_df <- data.frame(batch_sizes,epochs_run,training_accuracies,validation_accuracies)
performance_summary_df

```

##	batch_sizes	epochs_run	training_accuracies	validation_accuracies
## 1	50	10	0.7326531	0.7323357
## 2	50	15	0.7348032	0.7347164
## 3	50	20	0.7335641	0.7324207
## 4	50	40	0.7334912	0.7304651
## 5	50	60	0.7353134	0.7342913
## 6	50	100	0.7344024	0.7322506
## 7	100	10	0.7331268	0.7303801
## 8	100	15	0.7340379	0.7322506
## 9	100	20	0.7336735	0.7330159
## 10	100	40	0.7340744	0.7357367
## 11	100	60	0.7348032	0.7339512
## 12	100	100	0.7342566	0.7332710
## 13	150	10	0.7331268	0.7309753
## 14	150	15	0.7343295	0.7329309
## 15	150	20	0.7325802	0.7321656
## 16	150	40	0.7355685	0.7342913
## 17	150	60	0.7339286	0.7331010
## 18	150	100	0.7352770	0.7354817
## 19	200	10	0.7336006	0.7343763
## 20	200	15	0.7325437	0.7315704
## 21	200	20	0.7331997	0.7310603
## 22	200	40	0.7330904	0.7325908
## 23	200	60	0.7361881	0.7348865
## 24	200	100	0.7347303	0.7352266

Identified the number of Epochs and Batch Size which produced the highest validation accuracy. High validation accuracy is important to ensure that the model produces low variance estimate.

```
performance_summary_df$epochs_run[which.max(performance_summary_df$validation_accuracies)]
```

```
## [1] 40
```

```
performance_summary_df$batch_sizes[which.max(performance_summary_df$validation_accuracies)]
```

```
## [1] 100
```

Hence going ahead with the batch size and epochs that produce the maximum validation accuracy to fit the final ANN1 deep learning model.

When the validation accuracy is also as high as training accuracy, we can be assured that the model will produce low bias and low variance estimates.

```

dl_model_keras_ann1 <- model_keras_ann1('adam')

model_ann1_fit <- fit(verbose=0,
  object = dl_model_keras_ann1,
  x = as.matrix(x_train_disease_pred_tbl),
  y = y_train_vec,
  batch_size = performance_summary_df$batch_sizes[which.max(performance_summary_df$validation_accuracies)],
  epochs = performance_summary_df$epochs_run[which.max(performance_summary_df$validation_accuracies)],
  validation_split = 0.30
)

```

Making Predictions on the split Test Data using the final ANN1 DL model built using the tibble function and yardstick library

Accuracy and Recall are the 2 most important metrics because Recall helps in minimizing the False Negatives which are highly undesirable when doing disease prediction

```

yhat_keras_class_vec <- predict_classes(object = dl_model_keras_ann1, x = as.matrix(x_test_disease_pred_tbl)) %>% as.vector()
yhat_keras_prob_vec <- predict_proba(object = dl_model_keras_ann1, x = as.matrix(x_test_disease_pred_tbl)) %>% as.vector()
estimates_keras_tbl <- tibble::tibble(
  truth = as.factor(y_test_vec),
  estimate = as.factor(yhat_keras_class_vec),
  class_prob = yhat_keras_prob_vec
)
#estimates_keras_tbl

```

Evaluating the ANN1 model performance on the test dataset (Hold One out method), using the same set of metrics: Accuracy, Precision, Recall, F1 score and AUC score to evaluate the performance of the ANN1 Deep Learning model

Since the accuracy on the testing data is completely inline with the training accuracy we can be assured that the model is producing low bias and low variance estimate

```
library(yardstick)
options(yardstick.event_first = F)
ann1_accuracy <- estimates_keras_tbl %>% metrics(truth, estimate)
ACCURACY<- ann1_accuracy$.estimate[1]

ann1_auc <-estimates_keras_tbl %>% roc_auc(truth, class_prob)
AUC<-ann1_auc$.estimate[1]

ann1_precision <-estimates_keras_tbl %>% precision(truth, estimate)
PRECISION <- ann1_precision$.estimate[1]

ann1_recall <-estimates_keras_tbl %>% recall(truth, estimate)
RECALL<- ann1_recall$.estimate[1]

ann1_f1 <-estimates_keras_tbl %>% f_meas(truth, estimate)
F1 <- ann1_f1$.estimate[1]

data.frame(ACCURACY,AUC,PRECISION,RECALL,F1)
```

```
##      ACCURACY      AUC PRECISION    RECALL      F1
## 1 0.7333401 0.801648 0.7475183 0.7059303 0.7261293
```

```
#install.packages("pROC")
library(pROC)
```


Generated ROC curve and calculated Area Under Curve metric for the identified best performing ANN1 model

```
roc_curve <- roc(y_test_vec,yhat_keras_prob_vec)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
plot(roc_curve)
```

Obtained the area under the ROC curve for the best ANN1 model

```
paste("Area Under the ROC curve is :",auc(roc_curve))
```

```
## [1] "Area Under the ROC curve is : 0.801648020307732"
```

We can observe the improvement in model performance between ANN0 and ANN1

END OF BUILDING ANN1 Deep Learning Model (algorithm number 3)

2.4 Artificial Neural Network - 2 hidden layers (ANN2)

Using the same X train, X test dataset tables and Y train, Y test vectors

Building neural network architecture with 2 hidden layers (ANN2)

The number of nodes in the first hidden layer should be equal to the number of input attributes in the dataset, whereas the number of nodes in the second hidden layer can vary anywhere between 1 and two times of number of input attributes.

Using the relu activation function after the hidden layers as relu produces 0 as output for all negative inputs and the input value itself is produced as output if the input is > 0 . Hence, the relu activation function is more suitable for taking partial derivatives. The Sigmoid function itself is being used after the output layer to produce output between 0 and 1.

The Dropout layers are used to control overfitting by eliminating weights below a cutoff threshold in order to prevent low weights from overfitting the layers. Here we are removing weights below 10%.

Using the binary_crossentropy loss function since it is a binary classification problem

Written a function below, which returns a Neural Network with two hidden layers by taking in the optimizer as input parameter

```
library(keras)
library(tensorflow)

model_keras_ann2 <- function(opt) {

  model_keras_ann <- keras_model_sequential()
  model_keras_ann %>%
    ## Hidden Layer 1
    layer_dense(units = 16,
      kernel_initializer = "uniform",
      activation = "relu",
      input_shape = ncol(x_train_disease_pred_tbl)) %>%

    ## Hidden Layer 2
    layer_dense(units = 12,
      kernel_initializer = "uniform",
      activation = "relu") %>%
    layer_dropout(rate = 0.1)%>%

    ## Output Layer
    layer_dense(units = 1,
      kernel_initializer = "uniform",
      activation = "sigmoid") %>%
    compile(optimizer = opt,
      loss = "binary_crossentropy",
      metrics = c("accuracy"))
  )

  return(model_keras_ann)


}
```

Fitting the Neural Network with 2 hidden layers onto the train data

```
model_ann2_fit <- fit(verbose=0,
  object = model_keras_ann2('adam'),
  x = as.matrix(x_train_disease_pred_tbl),
  y = y_train_vec,
  batch_size = 50,
  epochs = 15,
  validation_split = 0.30
)
```

From the below plot also we can identify the suitable number of epochs based on the point where the curve begins to flatten

```
plot(model_ann2_fit)
```



Tuning the number of epochs and batch size parameters to see which combination produces the best validation accuracy using ADAM optimizer network by using the ANN2 function written above which returns a DL neural network with 2 hidden layers in every iteration of the loop, by taking the optimizer as the function input parameter.

The number of epochs & batch size should be chosen in a way that both training and validation accuracies are high thus ensuring that the model doesn't overfit by producing a high training accuracy and low validation accuracy

```

batch_size = c(50,100,150,200)
epochs = c(10,15,20,40,60,100)
model_ann2_tuned_fit <- c(0)
for (i in batch_size) {
  for (j in epochs) {

    model_fit <- fit(verbose=0,
object = model_keras_ann2('adam'),
x = as.matrix(x_train_disease_pred_tbl),
y = y_train_vec,
batch_size = i,
epochs = j,
validation_split = 0.30
)
    model_ann2_tuned_fit <- c(model_ann2_tuned_fit,model_fit)
  }
}

```

Storing the training and validation accuracies for different batch size and epoch values to compare the accuracies for each combination

```

batch_sizes <- c()
epochs_run <- c()
validation_accuracies <- c()
training_accuracies <- c()
for (i in seq(2,48,2))
{
  batch_sizes <- c(batch_sizes,model_ann2_tuned_fit[i]$params$batch_size)
  epochs_run <- c(epochs_run,model_ann2_tuned_fit[i]$params$epochs)
}

for (i in seq(3,49,2)) {
  validation_accuracies <- c(validation_accuracies,tail(model_ann2_tuned_fit[i]$metrics$val_accuracy,n=1))
  training_accuracies <- c(training_accuracies, tail(model_ann2_tuned_fit[i]$metrics$accuracy,n=1))
}

```

Table summarizing the model performances and their corresponding batch size and epoch hyperparameters

The batch_size value & epochs values with their corresponding training accuracy and validation accuracy can be seen in the below table

```

performance_summary_df <- data.frame(batch_sizes,epochs_run,training_accuracies,validation_accuracies)
performance_summary_df

```

##	batch_sizes	epochs_run	training accuracies	validation accuracies
## 1	50	10	0.7354956	0.7321656
## 2	50	15	0.7343295	0.7319105
## 3	50	20	0.7361152	0.7360769
## 4	50	40	0.7400146	0.7344614
## 5	50	60	0.7385204	0.7358218
## 6	50	100	0.7385204	0.7333560
## 7	100	10	0.7352405	0.7324207
## 8	100	15	0.7345481	0.7334411
## 9	100	20	0.7353498	0.7328458
## 10	100	40	0.7362610	0.7338662
## 11	100	60	0.7390306	0.7348015
## 12	100	100	0.7379009	0.7343763
## 13	150	10	0.7350219	0.7341213
## 14	150	15	0.7337463	0.7324207
## 15	150	20	0.7357872	0.7327608
## 16	150	40	0.7363338	0.7347164
## 17	150	60	0.7383747	0.7351416
## 18	150	100	0.7381560	0.7368421
## 19	200	10	0.7344388	0.7325057
## 20	200	15	0.7349125	0.7336111
## 21	200	20	0.7346939	0.7339512
## 22	200	40	0.7373178	0.7324207
## 23	200	60	0.7377186	0.7342913
## 24	200	100	0.7387755	0.7358218

Identified the number of Epochs and Batch Size which produced the highest validation accuracy

```
performance_summary_df$epochs_run[which.max(performance_summary_df$validation accuracies)]
```

```
## [1] 100
```

```
performance_summary_df$batch_sizes[which.max(performance_summary_df$validation accuracies)]
```

```
## [1] 150
```

Hence going ahead with the batch size and epochs that produce the maximum validation accuracy to fit the final ANN2 deep learning model

When the validation accuracy is also as high as training accuracy, we can be assured that the model will produce low bias and low variance estimates

```

dl_model_keras_ann2 <- model_keras_ann2('adam')

model_ann2_fit <- fit(verbose=0,
  object = dl_model_keras_ann2,
  x = as.matrix(x_train_disease_pred_tbl),
  y = y_train_vec,
  batch_size = performance_summary_df$batch_sizes[which.max(performance_summary_df$validation_accuracies)],
  epochs = performance_summary_df$epochs_run[which.max(performance_summary_df$validation_accuracies)],
  validation_split = 0.30
)

```

Making Predictions on the split Test Data using the final ANN2 DL model built using the tibble function and yardstick library

Accuracy and Recall are the 2 most important metrics because Recall helps in minimizing the False Negatives which are highly undesirable when doing disease prediction

```

yhat_keras_class_vec <- predict_classes(object = dl_model_keras_ann2, x = as.matrix(x_test_disease_pred_tbl)) %>% as.vector()
yhat_keras_prob_vec <- predict_proba(object = dl_model_keras_ann2, x = as.matrix(x_test_disease_pred_tbl)) %>% as.vector()
estimates_keras_tbl <- tibble::tibble(
  truth = as.factor(y_test_vec),
  estimate = as.factor(yhat_keras_class_vec),
  class_prob = yhat_keras_prob_vec
)
#estimates_keras_tbl

```

Evaluating the ANN2 model performance on the test dataset (Hold One out method), using the same set of metrics: Accuracy, Precision, Recall, F1 score and AUC score to evaluate the performance of the ANN2 Deep Learning model

Since the accuracy on the testing data is completely inline with the training accuracy we can be assured that the model is producing low bias and low variance estimate

```
library(yardstick)
options(yardstick.event_first = F)
ann2_accuracy <- estimates_keras_tbl %>% metrics(truth, estimate)
ACCURACY<- ann2_accuracy$.estimate[1]

ann2_auc <-estimates_keras_tbl %>% roc_auc(truth, class_prob)
AUC<-ann2_auc$.estimate[1]

ann2_precision <-estimates_keras_tbl %>% precision(truth, estimate)
PRECISION <- ann2_precision$.estimate[1]

ann2_recall <-estimates_keras_tbl %>% recall(truth, estimate)
RECALL<- ann2_recall$.estimate[1]

ann2_f1 <-estimates_keras_tbl %>% f_meas(truth, estimate)
F1 <- ann2_f1$.estimate[1]

data.frame(ACCURACY,AUC,PRECISION,RECALL,F1)
```

```
##      ACCURACY      AUC PRECISION  RECALL      F1
## 1 0.7398714 0.8055349 0.7531129 0.7148971 0.7335076
```

```
#install.packages("pROC")
library(pROC)
```


Generated ROC curve and calculated Area Under Curve metric for the identified best performing ANN2 model

```
roc_curve <- roc(y_test_vec,yhat_keras_prob_vec)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
plot(roc_curve)
```



Obtained the area under the ROC curve for the best ANN2 model

```
paste("Area Under the ROC curve is :",auc(roc_curve))
```

```
## [1] "Area Under the ROC curve is : 0.805534911746048"
```

END OF BUILDING ANN2 Deep Learning Model (algorithm number 4)

SECTION 3: Combination & Comparison of multiple Machine Learning Algorithms

3.1 Applying DECISION TREE INDUCTION

Generating Training and Validation datasets for decision tree with 70% & 30% splits respectively

```
library(caret)  
library(rpart)
```

```
## Warning: package 'rpart' was built under R version 3.6.2
```



```

set.seed(73)

disease_prediction_training$Disease <- as.factor(disease_prediction_training$Disease)
disease_prediction_training$Smoke <- as.factor(disease_prediction_training$Smoke)
disease_prediction_training$Alcohol <- as.factor(disease_prediction_training$Alcohol)
disease_prediction_training$Exercise <- as.factor(disease_prediction_training$Exercise)
disease_prediction_training$Height <- as.numeric(disease_prediction_training$Height)
disease_prediction_training$age_groups <- NULL
disease_prediction_training$bmi_groups <- NULL

train_index <- createDataPartition(disease_prediction_training$Disease, p = 0.7, list = FALSE)

disease_prediction_dt_tr_df <- disease_prediction_training[train_index, ]

disease_prediction_dt_test_df <- disease_prediction_training[-train_index, ]

```

Building a Decision Tree baseline model only by tuning the Complexity Parameter (CP) to produce best accuracy

```

#dt_model_disease <- train(Disease~., data = disease_prediction_dt_tr_df, method = "rpart",
# metric = "Accuracy",
# tuneLength = 8)

```

8 models were tuned and the model with lowest cp produced the highest accuracy, which was used as the final resulting model

```
print(dt_model_disease)
```

```

## CART
##
## 34301 samples
##    12 predictor
##    2 classes: '0', '1'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 34301, 34301, 34301, 34301, 34301, 34301, ...
## Resampling results across tuning parameters:
##
##   cp          Accuracy   Kappa
## 0.000971874  0.7330178  0.46615745
## 0.001166249  0.7327223  0.46554948
## 0.002157560  0.7312533  0.46256670
## 0.002449122  0.7307790  0.46163769
## 0.003032247  0.7294912  0.45908802
## 0.003323809  0.7293075  0.45871772
## 0.009709021  0.7234220  0.44699905
## 0.436235349  0.5414907  0.08659074
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.000971874.

```

Making Classifications on Validation Data using the baseline Decision Tree Model built above using the Training data

```
disease_dt_predict <- predict(dt_model_disease, newdata = disease_prediction_dt_test_df, na.action = na.omit, type = "raw")
```

Evaluating baseline decision tree Model Performance by calculating Accuracy Precision Recall & F1 score on the classification done on the Validation Data

```
disease_prediction_dt_test_df$predicted_disease_baseline <- disease_dt_predict
conf_matrix <- data.frame(table(disease_prediction_dt_test_df$Disease,disease_prediction_dt_test_df$predicted_disease_baseline))
colnames(conf_matrix)<- c('Actual class','Predicted Class','Count')
```

```
Accuracy_DT <- sum(conf_matrix$Count[conf_matrix$`Actual class`==conf_matrix$`Predicted Class`])/sum(conf_matrix$Count)
Precision_DT <- conf_matrix$Count[conf_matrix$`Actual class`=='1' & conf_matrix$`Predicted Class`=='1']/sum(conf_matrix$Count[conf_matrix$`Predicted Class`=='1'])
Recall_DT <- conf_matrix$Count[conf_matrix$`Actual class`=='1' & conf_matrix$`Predicted Class`=='1']/sum(conf_matrix$Count[conf_matrix$`Actual class`=='1'])
F1_score_DT <- (2*Precision_DT*Recall_DT)/(Precision_DT+Recall_DT)
```

```
paste("Accuracy of DT Algorithm in classifying Disease is :",Accuracy_DT)
```

```
## [1] "Accuracy of DT Algorithm in classifying Disease is : 0.72583168923056"
```

```
paste("Precision of DT Algorithm in classifying Disease is :",Precision_DT)
```

```
## [1] "Precision of DT Algorithm in classifying Disease is : 0.726182363363773"
```

```
paste("Recall of DT Algorithm in classifying Disease is :",Recall_DT)
```

```
## [1] "Recall of DT Algorithm in classifying Disease is : 0.724996598176623"
```

```
paste("F1 Score of DT Algorithm in classifying Disease is :",F1_score_DT)
```

```
## [1] "F1 Score of DT Algorithm in classifying Disease is : 0.725588996323029"
```

Introducing the Hyperparameters minbucket, minsplit and maxdepth alongwith CP for tuning the model

Minbucket is the minimum number of observations in any leaf node, **Minsplit** is the minimum number of observations that must exist in a node in order for a split to be attempted and **MaxDepth** is the maximum depth of any node of the final tree

Our aim is to tune these parameters to produce maximum accuracy, i.e. a model with minimum bias and at the same time the variance should not be too high by producing a good trade off between Bias and

Variance to avoid underfitting and overfitting.

Hence `minsplit`, `maxdepth` and `minbucket` should be adjusted so as to pre prune the model to produce a least complex model(to reduce overfitting and variance) with maximum accuracy (to minimize bias)

Creating a Grid with all combinations of 4 different values of `minsplit`, `maxdepth` and `minbucket`

```
library(tidyverse)
gs <- list(minsplit = c(20,15,10,5),
          maxdepth = c(30,25,20,15),
          minbucket = c(6,5,4,3)) %>%
  cross_d() # Convert to data frame grid
```

```
## Warning: `cross_d()` is deprecated; please use `cross_df()` instead.
```

Performing Grid Search with all combinations of different values of `minsplit`, `maxdepth` and `minbucket` created above and measuring the accuracy for each combination

```
#accuracy_values <- c()
#for (i in seq(1,nrow(gs))) {

# dt_model_disease_tuned <- train(Disease ~ ., data = disease_prediction_dt_tr_df,method = "rpart",metric = "Accuracy",
#                               tuneLength = 8,control = rpart.control(minsplit =as.numeric(gs[i,"minsplit"]), minbucket = as.numeric(gs[i,"minbucket"]), #maxdepth = as.numeric(gs[i,"maxdepth"])
#                               h""))
# accuracy_values<- c(accuracy_values,max(dt_model_disease_tuned$results$Accuracy))
#}
```

Displaying the Accuracy of 64 Decision Tree Models for different combinations of the hyperparameters `minsplit`, `maxdepth` and `minbucket` ordered in descending order by accuracy.

We can see that the first row in the below table with the below shown `minsplit`, `maxdepth` and `minbucket` values has the best accuracy and hence can be considered the best model.

But since the accuracy values are approximately equal, we will use a simpler model with lower `maxdepth` to avoid overfitting and produce low variance estimates. Based on Occam's Razor principle, we need to chose a simpler model, amongst the models that produce the same accuracy

```
gs$accuracy <- accuracy_values
gs[order(-gs$accuracy),]
```

```
## # A tibble: 64 x 4
##   minsplitted maxdepth minbucket accuracy
##   <dbl>      <dbl>      <dbl>      <dbl>
## 1         5        20         6      0.734
## 2        20        15         3      0.733
## 3        15        20         3      0.733
## 4        10        20         3      0.733
## 5        20        30         3      0.733
## 6        20        15         4      0.733
## 7        20        25         4      0.733
## 8        20        20         3      0.733
## 9        15        15         6      0.733
## 10       5         25         4      0.733
## # ... with 54 more rows
```

Creating a decision tree model with maxdepth 15 and the above obtained minsplitted & minbucket so as avoid overfitting and reduce estimate variance

```
# dt_model_disease_final_tuned <- train(Disease ~ ., data = disease_prediction_dt_tr_df, method
= "rpart", metric = "Accuracy",
#                               tuneLength = 8, control = rpart.control(minsplitted = gs$minsplitted[which.max
(gs$accuracy)], minbucket = #gs$minbucket[which.max(gs$accuracy)], maxdepth = 15))
# print(dt_model_disease_final_tuned)
```

Using the final best performing model that produced unbiased and low variance estimates to predict the classes (classify) on the validation dataset

```
disease_dt_prediction <- predict(dt_model_disease_final_tuned, newdata = disease_prediction_dt_t
est_df, na.action = na.omit, type = "raw")
```

Evaluating Model Performance by calculating Accuracy, Precision, Recall & F1 score on the classification done on the Test Data

The training accuracy and validation accuracy are in line, hence the model variance is low

```
disease_prediction_dt_test_df$predicted_disease_tuned <- disease_dt_prediction
conf_matrix <- data.frame(table(disease_prediction_dt_test_df$Disease, disease_prediction_dt_test
_df$predicted_disease_tuned))
colnames(conf_matrix) <- c('Actual class', 'Predicted Class', 'Count')

Accuracy_DT <- sum(conf_matrix$Count[conf_matrix$`Actual class` == conf_matrix$`Predicted Class
`])/sum(conf_matrix$Count)
Precision_DT <- conf_matrix$Count[conf_matrix$`Actual class` == '1' & conf_matrix$`Predicted Class
` == '1']/sum(conf_matrix$Count[conf_matrix$`Predicted Class` == '1'])
Recall_DT <- conf_matrix$Count[conf_matrix$`Actual class` == '1' & conf_matrix$`Predicted Class` ==
'1']/sum(conf_matrix$Count[conf_matrix$`Actual class` == '1'])
F1_score_DT <- (2*Precision_DT*Recall_DT)/(Precision_DT+Recall_DT)

paste("Accuracy of tuned DT Algorithm in classifying Disease :", Accuracy_DT)
```

```
## [1] "Accuracy of tuned DT Algorithm in classifying Disease : 0.72583168923056"
```

```
paste("Precision of tuned DT Algorithm in classifying Disease is :",Precision_DT)
```

```
## [1] "Precision of tuned DT Algorithm in classifying Disease is : 0.726182363363773"
```

```
paste("Recall of tuned DT Algorithm in classifying Disease is :",Recall_DT)
```

```
## [1] "Recall of tuned DT Algorithm in classifying Disease is : 0.724996598176623"
```

```
paste("F1 Score of tuned DT Algorithm in classifying Disease is :",F1_score_DT)
```

```
## [1] "F1 Score of tuned DT Algorithm in classifying Disease is : 0.725588996323029"
```

```
#install.packages("rattle")  
library(rattle)
```

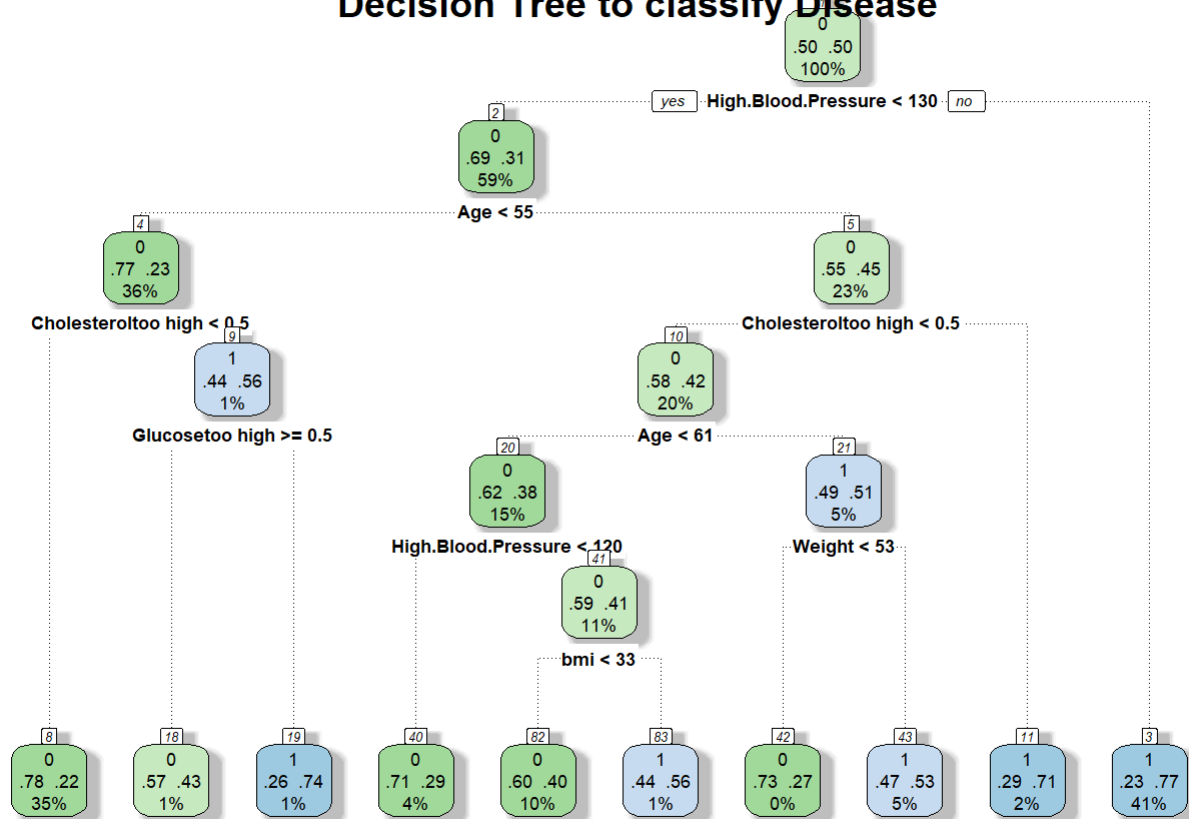
```
## Warning: package 'rattle' was built under R version 3.6.2
```

```
## Rattle: A free graphical interface for data science with R.  
## Version 5.3.0 Copyright (c) 2006-2018 Togaware Pty Ltd.  
## Type 'rattle()' to shake, rattle, and roll your data.
```

displaying the best performing Decision Tree model

```
fancyRpartPlot(dt_model_disease_final_tuned$finalModel, main = "Decision Tree to classify Disease")
```

Decision Tree to classify Disease



Rattle 2020-Apr-25 17:02:00 bhavi

ROC & AUROC for Decision Tree Induction

```
#install.packages("pROC")
library(pROC)
```


Generated ROC curve and calculated Area Under Curve metric for the identified best performing decision tree model

```
dt_pred_disease_prob <- predict(dt_model_disease_final_tuned, newdata = disease_prediction_dt_test_df, na.action = na.omit, type = "prob")
roc_curve <- roc(disease_prediction_dt_test_df$Disease, dt_pred_disease_prob[,1])
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
plot(roc_curve)
```



Obtained the area under the ROC curve

```
paste("Area Under the ROC curve is :",auc(roc_curve))
```

```
## [1] "Area Under the ROC curve is : 0.757847298396839"
```

Analyzing the importance of all the variables in the best performing Decision Tree model & comparing with the importance of variables as per other learning algorithms

1. The top 5 most important features as per Decision Tree are “High Blood Pressure”, “Low Blood Pressure”, “Age”, “Cholestrol high” & “cholestrol normal”

```
varImp(dt_model_disease_final_tuned)
```

```
## rpart variable importance
##
##               Overall
## High.Blood.Pressure 100.00000
## Low.Blood.Pressure  64.38748
## Age                 35.68229
## Cholesteroltoo high 32.23398
## Cholesterolnormal   31.92579
## bmi                 4.22820
## Exercise1           1.69968
## Weight              0.96418
## Glucosetoo high     0.64575
## Glucosenormal       0.52399
## Gendermale          0.09931
## Height              0.08340
## Alcohol1            0.00000
## Smoke1              0.00000
## `Cholesteroltoo high` 0.00000
## `Glucosetoo high`    0.00000
```

2. The top 5 most important variables as per Logistic Regression which have the lowest P-values and the highest absolute coefficients are, “High Blood Pressure”, “cholesterol_normal”, “Age”, “Cholesterol high” & “bmi”, which is mostly consistent with the top 5 most important variables as per Decision Tree. The only difference is that “bmi” is slightly more important than “low blood pressure” as per logistic regression, due to larger coefficient whereas the vice versa is true for Decision Tree.

```
#library(caret)
#logistic_summary_disease <- train(Disease ~ ., data = disease_prediction_training_df_tr_split,
  method = "glm", family = "binomial")
summary(logistic_summary_disease)
```



```
##
## Call:
## NULL
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.6029  -0.9217  -0.1818   0.9268   2.7994
##
## Coefficients: (2 not defined because of singularities)
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -4.28167    0.10239  -41.819  < 2e-16 ***
## Age           1.79840    0.06704   26.824  < 2e-16 ***
## High.Blood.Pressure 6.90232    0.15944   43.290  < 2e-16 ***
## Low.Blood.Pressure 1.30998    0.16306    8.034 9.45e-16 ***
## Smoke1       -0.14621    0.04897   -2.986  0.00283 **
## Alcohol1     -0.09309    0.05960   -1.562  0.11832
## Exercise1    -0.22280    0.03099   -7.188 6.55e-13 ***
## bmi          6.35296    0.69319    9.165  < 2e-16 ***
## Gender_female1 -0.01696    0.02781   -0.610  0.54184
## Cholesterol_high1 -0.66489    0.05958  -11.159  < 2e-16 ***
## Cholesterol_normal1 -1.08463    0.05061  -21.432  < 2e-16 ***
## `\\`Cholesterol_too high\\`1`      NA         NA         NA         NA
## Glucose_high1      0.31133    0.07311    4.258 2.06e-05 ***
## Glucose_normal1    0.30442    0.05605    5.431 5.60e-08 ***
## `\\`Glucose_too high\\`1`          NA         NA         NA         NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 47551  on 34300  degrees of freedom
## Residual deviance: 38386  on 34288  degrees of freedom
## AIC: 38412
##
## Number of Fisher Scoring iterations: 4
```

3. The top 5 most important variables as per Random Forest are “High Blood Pressure”, “Low Blood Pressure”, “Age”, “bmi” and “weight” as shown below. The top 3 are same as that of Decision Tree and only the 4th and 5th most important variables are different between Decision Tree and Random Forest.

```
varImp(tuned_model_rf)
```

```
## rf variable importance
##
##               Overall
## High.Blood.Pressure 100.00000
## Low.Blood.Pressure  54.69398
## Age                 29.27259
## bmi                 28.70573
## Weight              21.94487
## Height              16.66411
## Cholesterolnormal   12.30199
## Cholesteroltoo high 11.23895
## age_groups58 to 64  8.52774
## Glucosenormal       1.57926
## Exercise            1.54634
## age_groups53 to 58  1.17262
## Gendermale          0.99495
## Glucosetoo high     0.34094
## Smoke               0.33531
## age_groups48 to 53  0.02419
## Alcohol             0.00000
```

4. The top 5 most important variables as per Gradient Boosting Machine are “High Blood Pressure”, “Age”, “Cholestrol high”, “bmi” and “Cholestrol normal” as shown below. This is mostly consistent with the top 5 most important variables as per Decision Tree, with the only difference being that “Low Blood Pressure” is in the top 5 instead of “bmi” in case of Decision Tree, whereas the vice versa is true for Gradient Boosting Machine.

```
library(gbm)
```

```
## Warning: package 'gbm' was built under R version 3.6.3
```

```
## Loaded gbm 2.1.5
```

```
varImp(gbm_tuned)
```

```
## gbm variable importance
##
## Overall
## High.Blood.Pressure 100.0000
## Age 17.2772
## Cholesteroltoo high 6.4313
## bmi 3.6047
## Cholesterolnormal 3.2442
## Weight 2.3744
## Low.Blood.Pressure 2.2933
## Height 1.1410
## Exercise 0.9054
## Smoke 0.4171
## Glucosetoo high 0.3683
## Alcohol 0.1544
## Glucosenormal 0.1303
## Gendermale 0.0000
```

Creating Model Performance Master Table

The 2 most important performance evaluation metrics for this dataset are “ACCURACY” & “RECALL”. Since the dataset is balanced, with the number of people with disease and without disease being approximately equal, we can certainly use ‘accuracy’ to measure the model performance.

The next most important metric is “RECALL”. “Recall” is also extremely important when doing disease prediction, because we want to minimize false negatives, but without compromising much on ACCURACY. False Negatives are cases when the patient actually has the disease but our ML model falsely classifies the patient as not having the disease, which is highly undesirable and such cases have to be minimized by increasing our RECALL.

If we simply reduce the probability threshold to a very low value and classify everything as Positive, our False Negatives will be 0 and our RECALL may reach 100%, but then by doing so our ACCURACY will be very low, which will make it a bad model. Hence Recall has to be maximized without compromising on ACCURACY.

In our below master table, we can see that Accuracies of most of the models are almost inline with only a minor difference. Hence we will sort the table in decreasing order of their “RECALL” scores.

We can see the models with the highest RECALL score and also high AccURACY scores. We can observe that the Deep Learning Neural Network models (ANN0, ANN1, ANN2) and the Ensemble Models (GBM, RF) are one of the top performing models.

```

Algorithm <- c('KNN','Naive Bayes Classifier','Random Forest','Gradient Boosting Machine','Linear SVM','Non Linear SVM','Logistic Regression','Decision Tree','ANN0','ANN1','ANN2')
Hyperparameters <- c('K','Laplace','ntree,.mtry','depth,trees,shrinkage,minobs','C','sigma,C','alpha,lambda','minsplit,maxdepth,minbucket','batchsize,epochs','batchsize,epochs,layerdropout,units','batchsize,epochs,layerdropout,units')
Accuracy <- c(0.7284,0.7135,0.7341,0.7327,0.7288,0.7384,Accuracy_Logistic,Accuracy_DT,ann0_accuracy$.estimate[1],ann1_accuracy$.estimate[1],ann2_accuracy$.estimate[1])
Recall <- c(0.67914,0.6133,0.6871,0.7001,0.6117,0.6572,Recall_Logistic,Recall_DT,ann0_recall$.estimate[1],ann1_recall$.estimate[1],ann2_recall$.estimate[1])
F1_score <- c(0.7143,0.6834,0.7226,0.7253,0.6929,0.7152,F1_score_Logistic,F1_score_DT,ann0_f1$.estimate[1],ann1_f1$.estimate[1],ann2_f1$.estimate[1])
AUCs <- c(0.7918,0.7813,0.7913,0.8006,NA,NA,auc_logistic,auc(roc_curve),ann0_auc$.estimate[1],ann1_auc$.estimate[1],ann2_auc$.estimate[1])
ApproxRun_Time <- c('30 minutes','5 minutes','2 hours','1 hour','2 hours','2 hours','10 minutes','30 minutes','10 minutes','10 minutes','10 minutes')

master_table_df <- data.frame(Algorithm,Hyperparameters,Accuracy,Recall,ApproxRun_Time)

master_table_df <- master_table_df[order(-Recalls,-Accuracies),]

master_table_df

```

##	Algorithm	Hyperparameters	Accuracy	Recall	ApproxRun_Time
## 8	Decision Tree	minsplit,maxdepth,minbucket	0.7258317	0.7249966	30 m
## 11	ANN2	batchsize,epochs,layerdropout,units	0.7398714	0.7148971	10 m
## 10	ANN1	batchsize,epochs,layerdropout,units	0.7333401	0.7059303	10 m
## 4	Gradient Boosting Machine	depth,trees,shrinkage,minobs	0.7327000	0.7001000	1 hour
## 3	Random Forest	ntree,.mtry	0.7341000	0.6871000	2 hours
## 1	KNN	K	0.7284000	0.6791400	30 m
## 9	ANN0	batchsize,epochs	0.7272171	0.6788262	10 m
## 7	Logistic Regression	alpha,lambda	0.7273964	0.6640359	10 m
## 6	Non Linear SVM	sigma,C	0.7384000	0.6572000	2 hours
## 2	Naive Bayes Classifier	Laplace	0.7135000	0.6133000	5 m
## 5	Linear SVM	C	0.7288000	0.6117000	2 hours

END of Combination & Comparison of multiple algorithms

SECTION 4: PREDICTION and INTERPRETATION

Importing the TEST Dataset to apply all the built models on the test dataset to predict if each person in the testing dataset has the disease

```
setwd("D:/SYR ADS/Sem 2/IST_707_Data_Analytics/HW4")  
  
getwd
```

```
## function ()  
## .Internal(getwd())  
## <bytecode: 0x000002478470eac8>  
## <environment: namespace:base>
```

```
disease_prediction_testing <- read.csv("Disease Prediction Testing.csv")
```

VIEWING THE STRUCTURE and SUMMARY STATISTICS of the Data and checking for missing values

```
str(disease_prediction_testing)
```

```
## 'data.frame':    21000 obs. of  12 variables:  
## $ ID              : int  0 1 2 3 4 5 6 7 8 9 ...  
## $ Age              : int  44 41 63 55 55 58 45 52 58 52 ...  
## $ Gender           : Factor w/ 2 levels "female","male": 1 1 2 1 1 1 1 1 2 1 ...  
## $ Height           : int  160 169 168 158 167 162 161 149 168 165 ...  
## $ Weight           : num  59 74 84 108 67 95 68 85 64 92 ...  
## $ High.Blood.Pressure: int  100 120 120 160 120 130 120 160 140 150 ...  
## $ Low.Blood.Pressure : int  80 70 80 100 80 70 70 90 90 100 ...  
## $ Cholesterol       : Factor w/ 3 levels "high","normal",...: 1 2 2 2 2 2 2 2 2 2 ...  
## $ Glucose           : Factor w/ 3 levels "high","normal",...: 2 2 1 2 2 2 2 2 2 2 ...  
## $ Smoke             : int  0 0 0 0 0 0 0 0 0 0 ...  
## $ Alcohol           : int  0 0 0 0 0 0 0 0 0 0 ...  
## $ Exercise          : int  1 1 1 0 1 1 1 1 1 1 ...
```

```
summary(disease_prediction_testing)
```

```
##      ID      Age      Gender      Height      Weight      High.Blood.P
ressure Low.Blood.Pressure  Cholesterol      Glucose
## Min.    :    0  Min.    :29.00  female:13667  Min.    : 64.0  Min.    : 21.00  Min.    :  1
0.0      Min.    : -70.00  high    : 2844  high    : 1563
## 1st Qu.: 5250  1st Qu.:48.00  male   : 7333  1st Qu.:159.0  1st Qu.: 65.00  1st Qu.: 12
0.0      1st Qu.: 80.00  normal :15709  normal :17827
## Median :10500  Median :53.00                      Median :165.0  Median : 72.00  Median : 12
0.0      Median : 80.00  too high: 2447  too high: 1610
## Mean    :10500  Mean    :52.81                      Mean    :164.3  Mean    : 74.24  Mean    : 12
9.1      Mean    : 95.96
## 3rd Qu.:15749  3rd Qu.:58.00                      3rd Qu.:170.0  3rd Qu.: 82.00  3rd Qu.: 14
0.0      3rd Qu.: 90.00
## Max.    :20999  Max.    :64.00                      Max.    :250.0  Max.    :183.00  Max.    :1602
0.0      Max.    :8500.00
##      Smoke      Alcohol      Exercise
## Min.    :0.00000  Min.    :0.00000  Min.    :0.000
## 1st Qu.:0.00000  1st Qu.:0.00000  1st Qu.:1.000
## Median :0.00000  Median :0.00000  Median :1.000
## Mean    :0.08781  Mean    :0.05267  Mean    :0.805
## 3rd Qu.:0.00000  3rd Qu.:0.00000  3rd Qu.:1.000
## Max.    :1.00000  Max.    :1.00000  Max.    :1.000
```

DATA PREPARATION of Testing Data

From the structure and summary of the data we can observe that the Min and Max values of the columns Low Blood Pressure and High Blood Pressure are not practically possible values and hence they are noise/outliers which need to be treated. Hence these columns need to be winsorized.

Winsorization is a data treatment process where the extreme outlier values are replaced with less extreme values which are practically possible

The min & max of low BP (diastolic BP) fall in the range of 45 to 140 and hence any value that is less than 45 is replaced with 45 and any value greater than 140 is replaced with 140

```
disease_prediction_testing$Low.Blood.Pressure[disease_prediction_testing$Low.Blood.Pressure<45]
<- 45

disease_prediction_testing$Low.Blood.Pressure[disease_prediction_testing$Low.Blood.Pressure>140]
<- 140
```

Verifying that the Min & Max values of Low Blood Pressure (Diastolic BP) are in the correct practically permissible range and the outliers have been eliminated

```
summary(disease_prediction_testing$Low.Blood.Pressure)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   45.00  80.00   80.00   82.13  90.00  140.00
```

The possible values for the min & max of High BP (Systolic BP) fall in the range of 70 to 200 and hence any value that is less than 70 is replaced with 70 and any value greater than 200 is replaced with 200

```
disease_prediction_testing$High.Blood.Pressure[disease_prediction_testing$High.Blood.Pressure<70] <- 70

disease_prediction_testing$High.Blood.Pressure[disease_prediction_testing$High.Blood.Pressure>200] <- 200
```

Verifying that the Min & Max values of High Blood Pressure (Systolic BP) are in the correct practically permissible range and the outliers have been eliminated

```
summary(disease_prediction_testing$High.Blood.Pressure)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      70.0   120.0   120.0   126.8   140.0   200.0
```

There are 84 instances where Low BP is > high BP even after winsorizing which needs to be treated by swapping the values

```
length(disease_prediction_testing[disease_prediction_testing$Low.Blood.Pressure>disease_prediction_testing$High.Blood.Pressure,1])
```

```
## [1] 84
```

Swapping the values wherever Low BP > High BP which is not permissible

```
low_bp_values <- disease_prediction_testing$Low.Blood.Pressure[disease_prediction_testing$Low.Blood.Pressure>disease_prediction_testing$High.Blood.Pressure]
```

```
high_bp_values <- disease_prediction_testing$High.Blood.Pressure[disease_prediction_testing$Low.Blood.Pressure>disease_prediction_testing$High.Blood.Pressure]
```

```
disease_prediction_testing$Low.Blood.Pressure[disease_prediction_testing$Low.Blood.Pressure>disease_prediction_testing$High.Blood.Pressure] <- high_bp_values
```

```
disease_prediction_testing$High.Blood.Pressure[disease_prediction_testing$Low.Blood.Pressure>disease_prediction_testing$High.Blood.Pressure] <- low_bp_values
```

Verifying that there are no instances with low bp values > high bp values

```
length(disease_prediction_testing[disease_prediction_testing$Low.Blood.Pressure>disease_prediction_testing$High.Blood.Pressure,1])
```

```
## [1] 0
```

Any value that is less than 28.9 are replaced with 28.9

```
disease_prediction_testing$Weight[disease_prediction_testing$Weight<28.9] <- 28.9
```

Verifying that the Min & Max values of Weight are in the correct practically permissible range and the outliers have been eliminated

```
summary(disease_prediction_testing$Weight)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  28.90   65.00   72.00   74.24   82.00  183.00
```

DATA PREPARATION for LOGISTIC REGRESSION:

For LOGISTIC REGRESSION, we require all the categorical variables to be converted to dummies by performing one hot encoding and also all the numeric variables have to be normalized so that all the columns have values in the range of 0 to 1.

Creating Dummy Variables out of all categorical variables by performing one hot encoding and normalizing all numeric columns using Min Max scaler

```
normalize <- function(x) {
  return ((x - min(x)) / (max(x) - min(x)))
}
library(fastDummies)
disease_prediction_testing_df <- disease_prediction_testing
disease_prediction_testing_df <- fastDummies::dummy_cols(disease_prediction_testing_df, select_columns=c('Gender', 'Cholesterol', 'Glucose'))

disease_prediction_testing_df$bmi <- disease_prediction_testing_df$Weight/((disease_prediction_testing_df$Height/100)*(disease_prediction_testing_df$Height/100))
disease_prediction_testing_df$Age <- normalize(disease_prediction_testing_df$Age)
disease_prediction_testing_df$Height <- normalize(disease_prediction_testing_df$Height)
disease_prediction_testing_df$Weight <- normalize(disease_prediction_testing_df$Weight)
disease_prediction_testing_df$bmi <- normalize(disease_prediction_testing_df$bmi)
disease_prediction_testing_df$Low.Blood.Pressure <- normalize(disease_prediction_testing_df$Low.Blood.Pressure)
disease_prediction_testing_df$High.Blood.Pressure <- normalize(disease_prediction_testing_df$High.Blood.Pressure)

i <- sapply(disease_prediction_testing_df, is.integer)
disease_prediction_testing_df[,i] <- lapply(disease_prediction_testing_df[,i], as.factor)

disease_prediction_testing_df$Gender <- NULL
disease_prediction_testing_df$Cholesterol <- NULL
disease_prediction_testing_df$Glucose <- NULL
disease_prediction_testing_df$age_groups <- NULL
disease_prediction_testing_df$ID <- NULL
disease_prediction_testing_df$bmi_groups <- NULL
disease_prediction_testing_df$Gender_male <- NULL
disease_prediction_testing_df$ID <- NULL
```

1. Obtaining Predicted values for disease on TEST data using the final LOGISTIC REGRESSION model**

```
LR <- predict(tuned_model_logistic, newdata = disease_prediction_testing_df)
```

PREPARING DATA for ANN

step_center() to mean-center the data & step_scale() to scale the data**The last step is to prepare the recipe with the prep() function**

```
library(recipes)
```

```
## Warning: package 'recipes' was built under R version 3.6.3
```

```
##  
## Attaching package: 'recipes'
```

```
## The following object is masked from 'package:stringr':  
##  
##      fixed
```

```
## The following object is masked from 'package:stats':  
##  
##      step
```

```
disease_rec_obj_test <- recipes::recipe(x = disease_prediction_testing_df)%>%  
  recipes::step_center(Age,Height,Weight,High.Blood.Pressure,Low.Blood.Pressure,bmi)%>%  
  recipes::step_scale(Age,Height,Weight,High.Blood.Pressure,Low.Blood.Pressure,bmi)%>%  
  recipes::prep(data = disease_prediction_testing_df)
```

Applying the “recipe” to data set with the bake() function which processes the data following our recipe steps

```
disease_prediction_testing_baked_df <- recipes::bake(disease_rec_obj_test, new_data = disease_prediction_testing_df)  
  
i <- sapply(disease_prediction_testing_baked_df, is.factor)  
disease_prediction_testing_baked_df[,i] <- lapply(disease_prediction_testing_baked_df[i], as.character)  
  
j <- sapply(disease_prediction_testing_baked_df, is.character)  
disease_prediction_testing_baked_df[,j] <- lapply(disease_prediction_testing_baked_df[j], as.integer)
```

2. Obtaining Predicted values for disease on TEST data using the final ANN0, ANN1 & ANN2 models

```
library(yardstick)
library(keras)
library(tensorflow)
ANN2 <- predict_classes(object = dl_model_keras_ann2, x = as.matrix(disease_prediction_testing_b
aked_df)) %>% as.vector()

ANN1 <- predict_classes(object = dl_model_keras_ann1, x = as.matrix(disease_prediction_testing_b
aked_df)) %>% as.vector()

ANN0 <- predict_classes(object = dl_model_keras_ann0, x = as.matrix(disease_prediction_testing_b
aked_df)) %>% as.vector()
```

Preparing Data for Decision Tree

```
disease_prediction_testing_dt <- disease_prediction_testing
disease_prediction_testing_dt$bmi <- disease_prediction_testing_dt$Weight/((disease_prediction_t
esting_dt$Height/100)*(disease_prediction_testing_dt$Height/100))
disease_prediction_testing_dt$ID <- NULL
disease_prediction_testing_dt$Smoke <- as.factor(disease_prediction_testing_dt$Smoke)
disease_prediction_testing_dt$Alcohol <- as.factor(disease_prediction_testing_dt$Alcohol)
disease_prediction_testing_dt$Exercise <- as.factor(disease_prediction_testing_dt$Exercise)
disease_prediction_testing_dt$Height <- as.numeric(disease_prediction_testing_dt$Height)
```

3. Obtaining Predicted values for disease on TEST data using the DECISION TREE model

```
DT <- predict(dt_model_disease_final_tuned, newdata = disease_prediction_testing_dt, na.action =
na.omit, type = "raw")
```

Writing the final predictions to a CSV

```
disease_prediction_testing$DT <- DT
disease_prediction_testing$LR <- LR
disease_prediction_testing$ANN0 <- ANN0
disease_prediction_testing$ANN1 <- ANN1
disease_prediction_testing$ANN2 <- ANN2

final_predictions <- disease_prediction_testing[,c('ID','DT','LR','ANN0','ANN1','ANN2')]
write.csv(final_predictions,"HW_04_Kumar_Bhavishh_Predictions.csv", row.names = FALSE)
```

D. CONCLUSION:

From the above model building, tuning and evaluation we can understand the importance of tuning hyperparameters through a grid search to improve the model performance. We also learned the importance of K fold cross validation and hold one out techniques to ensure that the model performs well not only on the training data but also on the validation data which it has not seen before. From the above 5 algorithms we can observe that the Deep Learning Models like Artificial Neural Networks with 0 or more hidden layers are one of the best performing models. For ANN, we can tune the batch size and epochs along with number of nodes and dropout rate in order to produce low bias and low variance estimates such that the model performs well not only on training data, but also on validation data.